

# Automatic Transactional Systems: Project

Authors: Adam Foster and Maciek Staniszewski

# Introduction

Forecasting stock prices has been a well-known challenge for financial analysts. As part of our final project we forecasted returns and close prices of a time series of a chosen stock (Johnson & Johnson). We developed two models: AR-GARCH and LSTM for return/price prediction. We built a long-short trading setup based on signals from these predictions. We assessed performance of each of the models using Sharpe ratio on the out-of-sample period, which consisted the last 10% of the time series.

Global variables were defined to be used across both trading strategies. They were initialized with capital letters.

```
In [ ]: TRANSACTION_COSTS = 0.0005 # cost is 5bps of trade value
```

## Literature Review

AR (Autoregressive) models are frequently used across statistics; also in trading strategies, as they provide a framework for explaining variables with lagged iterations of the variables. GARCH (Generalized Autoregressive Conditional Heteroskedasticity) models are widely used in trading strategies for their ability to model and forecast volatility in financial time series, capturing time-varying volatility clustering effectively. LSTM (Long Short-Term Memory) models, a type of recurrent neural network, excel in handling sequential data and capturing long-term dependencies, making them suitable for predicting stock prices and trends based on historical data. Recent studies highlight that while GARCH models are robust in volatility prediction, LSTM models often outperform in price prediction due to their non-linear processing capabilities. Combining both models can leverage GARCH's strength in volatility estimation and LSTM's prowess in trend prediction, potentially enhancing trading strategy performance.

Data: AR-GARCH

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
from statsmodels.tsa.stattools import acf
from statsmodels.tsa.stattools import pacf
from statsmodels.graphics.tsaplots import plot_acf
from statsmodels.stats.descriptivestats import describe
from statsmodels.stats.diagnostic import het_arch
from statsmodels.stats.diagnostic import acorr_ljungbox
import scipy as scipy
import arch
```

Johnson & Johnson, a pharmaceutical company listed on the NYSE, was chosen for this project. It is a well established company contributing to the S&P500 index with no long-term directional movement in price over the last several years. It is a good candidate for testing the true power of autoregressive and machine learning models and displays high liquidity.

Instead of 6 years of close prices, 2 months of 5-minute prices from 19/04/2024-07/06/2024 were used. This translated to 2574 observations, more than the required 1512 (252 trading days x 6 years alternative).

As a fallback, missing values were populated with last available observations. Log returns were calculated on close prices and scaled by 1000 for model calibration to prevent obtaining inaccurate and spurious results. These were converted back for calculating P&L.

```
In [ ]: jnj_raw = pd.read_csv('../jnj.us.txt')[["<DATE>", "<TIME>", "<OPEN>", "<CLOSE>"]].rename(columns={'<DATE>': 'DATE', '<TIME>': 'TIME', '<OPEN>': 'OPEN', '<CLOSE>': 'CLOSE'})

jnj_raw
```

```
Out[ ]:
```

	Date	Time	Open	Close
0	20240419	153000	146.150	144.8000
1	20240419	153500	144.800	145.4200
2	20240419	154000	145.440	145.6200
3	20240419	154500	145.660	145.9858
4	20240419	155000	145.975	146.1800
...	...	...	...	...
2569	20240607	213500	147.250	147.1700
2570	20240607	214000	147.170	147.0500
2571	20240607	214500	147.050	146.9300
2572	20240607	215000	146.930	147.1700
2573	20240607	215500	147.175	147.0800

2574 rows × 4 columns

```
In [ ]: jnj_raw['DateStr'] = jnj_raw['Date'].astype(str)
jnj_raw['TimeStr'] = jnj_raw['Time'].astype(str)

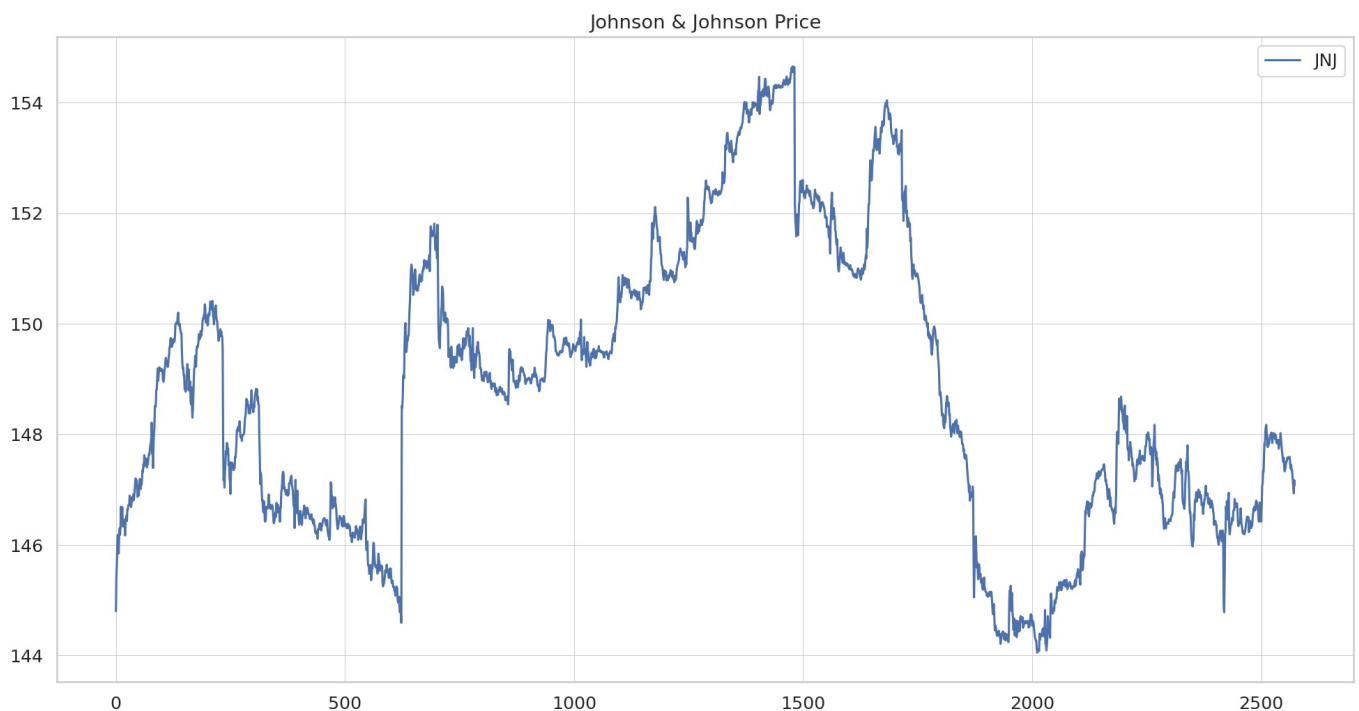
jnj_raw['Date'] = jnj_raw['DateStr'].apply(lambda x: f"{x[:4]}-{x[4:6]}-{x[6:]}")
jnj_raw['Time'] = jnj_raw['TimeStr'].apply(lambda x: f"{x[:2]}:{x[2:4]}:{x[4:]}")

jnj_raw['Datetime'] = pd.to_datetime(jnj_raw['DateStr'] + ' ' + jnj_raw['TimeStr'])
jnj_raw = jnj_raw.drop(columns = ['Open'])
jnj_raw = jnj_raw.rename(columns = {'Close': 'JNJ'})
```

```
In [ ]: plt.figure(figsize=(16, 8), dpi=150)

jnj_raw['JNJ'].plot(label='JNJ')
plt.title('Johnson & Johnson Price')
plt.legend()
```

```
Out[ ]: <matplotlib.legend.Legend at 0x7fe85179dd50>
```



```
In [ ]: returns = jnj_raw.fillna(method='ffill')
tickers = ['JNJ']

for ticker in tickers:
    returns['Ret_' + ticker] = (np.log(returns[ticker]) - np.log(returns[ticker].shift(1))) * 1000 # series sca

returns = returns.dropna() # only nulls left are the starting returns; will also have an issue with negative pr
returns.head()
```

/tmp/ipykernel\_6274/2988824415.py:1: FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.

```
returns = jnj_raw.fillna(method='ffill')
```

Out [ ]:

	Date	Time	JNJ	DateStr	TimeStr	Datetime	Ret_JNJ
1	2024-04-19	15:35:00	145.4200	20240419	153500	2024-04-19 15:35:00	4.272627
2	2024-04-19	15:40:00	145.6200	20240419	154000	2024-04-19 15:40:00	1.374382
3	2024-04-19	15:45:00	145.9858	20240419	154500	2024-04-19 15:45:00	2.508868
4	2024-04-19	15:50:00	146.1800	20240419	155000	2024-04-19 15:50:00	1.329382
5	2024-04-19	15:55:00	146.0800	20240419	155500	2024-04-19 15:55:00	-0.684322

## Statistics

In [ ]:

```
class data_extraction:
    def __init__(self, data, asset):
        self.data = data
        self.asset = asset
        self.series = self.extract_series()

    def extract_series(self):
        return self.data[self.asset]
```

In [ ]:

```
class analysis(data_extraction):
    def __init__(self, data, asset):
        data_extraction.__init__(self, data, asset)
        self.x = self.norm_dist()

    def acf_log_ret(self):
        plot_acf(self.series, lags=30, title='ACF Log Returns ' + self.asset)
        plt.show()

    def acf_sq_log_ret(self):
        plot_acf(self.series**2, lags=30, title='ACF Squared Log Returns ' + self.asset)
        plt.show()

    def describe(self):
        return describe(self.series)

    def norm_dist(self):
        mu, sigma = np.mean(self.series), np.std(self.series)
        x = np.random.normal(mu, sigma, 1000)
        x = pd.Series(x, name='Normal Distribution')
        return x

    def asset_hist_v_norm(self):
        fig, ax1 = plt.subplots()
        ax2 = ax1.twinx()
        ax2.grid(False)
        ax1.hist(self.series, bins=30)
        g1 = sns.kdeplot(self.x, ax=ax2, color='r')
        g1.set(ylabel=None)
        g1.set(yticklabels=[])
        plt.title('Histogram vs Normal Distribution ' + self.asset);

    def asset_dist_v_norm(self):
        sns.set_style('whitegrid')
        series_norm = pd.concat([self.series, self.x], axis=1)
        sns.kdeplot(data=series_norm, bw_method=0.5)
        plt.title('Distribution vs Normal Distribution ' + self.asset);

    def arch_test(self):
        return het_arch(self.series)

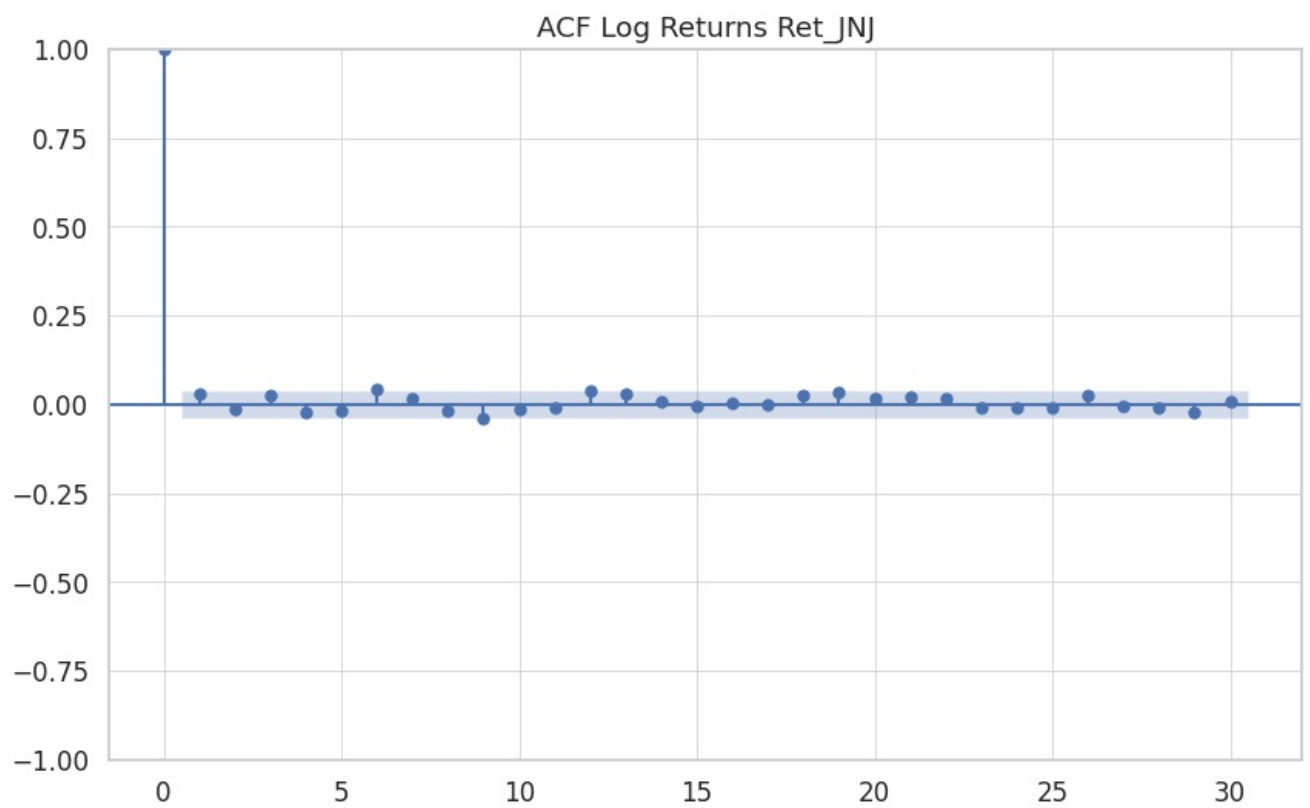
    def arch_test_lag5(self):
        return het_arch(self.series, nlags=5)
```

In [ ]:

```
analysis_jnj = analysis(returns, 'Ret_JNJ')
```

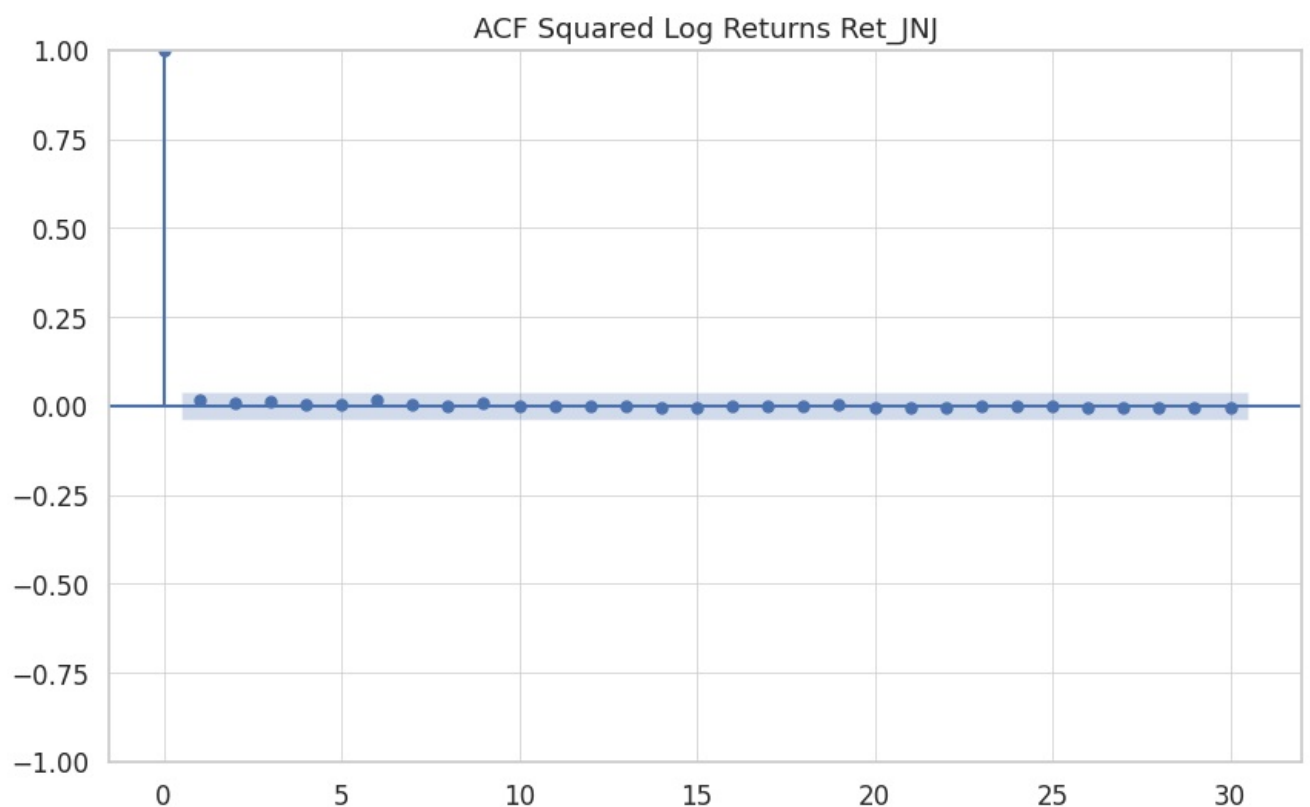
In [ ]:

```
analysis_jnj.acf_log_ret()
```



ACF of log returns indicates some AR/MA property among returns with lags 1, 4, 6, 6 being or almost being statistically significant at 5% confidence level.

```
In [ ]: analysis_jnj.acf_sq_log_ret()
```



ACF of squared log returns indicates low likelihood of AR property among squared returns for (G)ARCH with lags not being statistically significant.

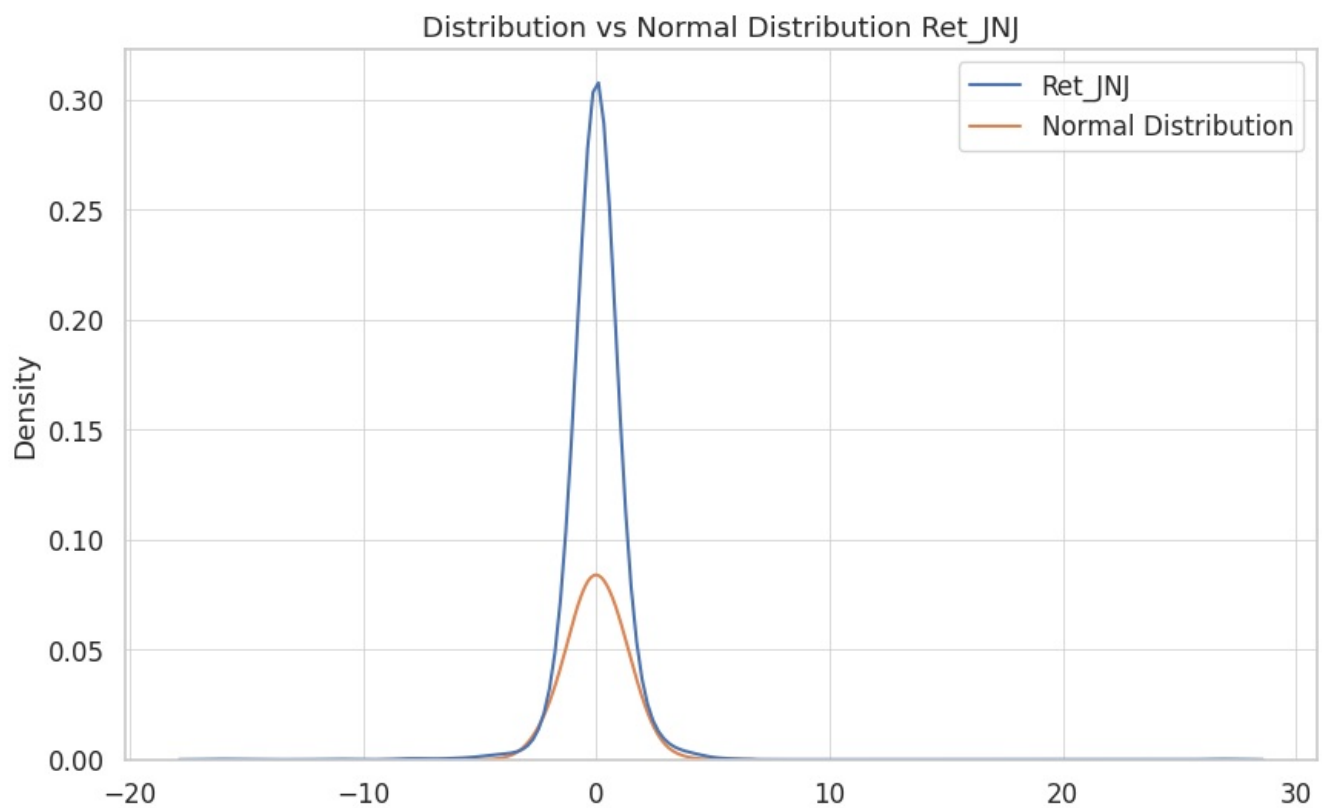
```
In [ ]: analysis_jnj.describe()
```

Out[ ]: Ret\_JNJ

nobs	2.573000e+03
missing	0.000000e+00
mean	6.071969e-03
std_err	2.346566e-02
upper_ci	5.206382e-02
lower_ci	-3.991988e-02
std	1.190290e+00
iqr	8.986146e-01
iqr_normal	6.661440e-01
mad	6.749249e-01
mad_normal	8.458929e-01
coef_var	1.960303e+02
range	4.288533e+01
max	2.678381e+01
min	-1.610152e+01
skew	2.789157e+00
kurtosis	1.216685e+02
jarque_bera	1.513068e+06
jarque_bera_pval	0.000000e+00
mode	0.000000e+00
mode_freq	2.409639e-02
median	0.000000e+00
1%	-3.127120e+00
5%	-1.379044e+00
10%	-9.634844e-01
25%	-4.444967e-01
50%	0.000000e+00
75%	4.541179e-01
90%	9.688703e-01
95%	1.474602e+00
99%	3.028375e+00

- Positive skewness: larger positive returns than negative ones are prevalent, which is unusual for equities
- Very high kurtosis: fluctuations in returns resulting in much fatter tails than expected of a normal distribution; somewhat expected given the frequency of the data
- jarque\_bera\_pval close to zero: H0 normality rejected

In [ ]: analysis\_jnj.asset\_dist\_v\_norm()



```
In [ ]: analysis_jnj.arch_test()
```

```
Out[ ]: (2.035973237564, 0.9960569359573243, 0.202884681235908, 0.9960952426298892)
```

```
In [ ]: analysis_jnj.arch_test_lag5()
```

```
Out[ ]: (1.0866108654935305,
0.9552819430219457,
0.2169061916290882,
0.9554283519084472)
```

- H0 no arch effects not rejected (large p-value): likely no arch effects in data

There were some mixed results regarding persistence of returns of Johnson & Johnson equity price. Whilst there were no arch effects, there was some indication of autoregression in the mean equation and an opportunity to exploit the fat tails of the leptokurtic distribution.

## Methodology: AR-GARCH

### Model

The first model was AR-GARCH. The AR (autoregressive) component defines the mean equation and regresses the variable on lags of itself with the lags specified denoting the persistence. The GARCH (generalised autoregressive conditional heteroskedasticity) component defines the variance equation and regresses variance on lagged variance and lagged squared error terms.

AR orders of 1, 4, 6 and 9 were selected for different specifications of the model in line with the ACF plot from the Statistics section.

GARCH(p,q) variations of (1,1), (1,2) and (2,1) were tested in the absence of statistically significant squared log returns and general lack of model improvement beyond two lags.

The model variations were run on 90% of the dataset as the train set. The remaining 10% was left for the test set. A rolling window approach was incorporated, whereby the parameters calibrated during training were applied to datetimes in the test set sequentially, forecasting only the next period return and volatility at a time.

```
In [ ]: class garch_base(data_extraction):
def __init__(self, approach, proportion, p, q, mean, data, asset, lags=0, o=0):
    data_extraction.__init__(self, data, asset)
    self.approach = approach
    self.proportion = proportion
    self.p = p
    self.o = o
    self.q = q
    self.mean = mean
    self.lags = lags
```

```

        self.train = self.series.iloc[:int(self.proportion*len(self.series))]
        self.test = self.series.iloc[int(self.proportion*len(self.series)):]
#         self.last_train_date = datetime.strptime(self.train.index[-1], '%Y-%m-%d')
        self.last_train_date = self.train.index[-1]
        self.res_garch = self.fit()

    def specs(self):
        return 'ar' + str(self.lags) + 'garch' + str(self.p) + str(self.q)

    def train_out(self):
        return self.train

    def fit(self):
        if self.approach == 'fixed':
            garch = arch.arch_model(self.train, mean=self.mean, lags=self.lags, vol='GARCH', p=self.p, o=self.o)
            res_garch = garch.fit()
        elif self.approach == 'roll1d':
            garch = arch.arch_model(self.series, mean=self.mean, lags=self.lags, vol='GARCH', p=self.p, o=self.o)
            res_garch = garch.fit(last_obs=self.last_train_date)

        return res_garch

    def summary(self):
        res_garch_summary = self.res_garch.summary()

        return res_garch_summary

    def plot(self):
        fig = self.res_garch.plot()
        plt.show()

    def autocorr(self):
        lb = acorr_ljungbox(self.res_garch.std_resid**2, [10, 15, 20])

        return lb

    def info_crit(self):
        ic_out = {'AIC': [self.res_garch.aic], 'BIC': [self.res_garch.bic]}
        ic = pd.DataFrame(data=ic_out)

        return ic

    def acf_stan(self):
        plot_acf(self.res_garch.std_resid, lags=30, title='ACF of Standardized Residuals')
        plt.show()

    def acf_sq_stan(self):
        plot_acf(self.res_garch.std_resid**2, lags=30, title='ACF of Squared Standardized Residuals')
        plt.show()

class garch(garch_base):
    def __init__(self, approach, proportion, p, q, mean, data, asset, lags=0, o=0):
        garch_base.__init__(self, approach, proportion, p, q, mean, data, asset, lags, o)
        self.forecast = self.forecast()
        self.exclusion_dates = self.get_exclusion_dates()
        self.forecast_mean_values = self.forecast_mean_values()
        self.forecast_std_values = self.forecast_std_values()
        self.forecast_vs_test = self.forecast_vs_test()

    def forecast(self):
        forecast = self.res_garch.forecast(horizon=len(self.test))

        return forecast

#     def forecast_out(self):
#         return self.forecast

    def get_exclusion_dates(self):
#         return [date for date in self.forecast.mean.index.strftime('%Y-%m-%d').tolist() if date <= self.last_train_date]
        return [i for i in self.forecast.mean.index.tolist() if i <= self.last_train_date]

    def forecast_mean_values(self):
        if self.approach == 'fixed':
            forecast_mean_values = self.forecast.mean.iloc[0]
        elif self.approach == 'roll1d':
            forecast_mean_values = self.forecast.mean.drop(self.exclusion_dates)
            forecast_mean_values = forecast_mean_values.iloc[:,0]

        return forecast_mean_values

    def forecast_mean(self):
#         forecast_mean = self.forecast.mean
#         forecast_mean.iloc[0].plot()
        self.forecast_mean_values.plot()

```

```

def forecast_std_values(self):
    forecast_vol_raw = self.forecast.variance**0.5

    if self.approach == 'fixed':
        forecast_std_values = forecast_vol_raw.iloc[0]
    elif self.approach == 'roll1d':
        forecast_std_values = forecast_vol_raw.drop(self.exclusion_dates)
        forecast_std_values = forecast_std_values.iloc[:,0]

    return forecast_std_values
#
#     std3 = std2.transpose()
#     for i, col in enumerate(std3.columns):
#         std3[col] = std3[col].shift(i)
#
#     std3
#     std3.plot()

def forecast_std(self):
#     forecast_stdev = self.forecast.variance**0.5
#     forecast_stdev.iloc[0].plot()
    self.forecast_std_values.plot()

def forecast_vs_test(self):
#     if self.approach == 'fixed':
#         forecast_mean_values = self.forecast.mean.iloc[0].values
#     elif self.approach == 'roll1d':
#         exclusion_dates = [date for date in self.forecast.mean.index.strftime('%Y-%m-%d').tolist() if date
#         forecast_mean_values = self.forecast.mean.drop(exclusion_dates)
#         forecast_mean_values = forecast_mean_values.iloc[:,0].values

    forecast_mean_values = self.forecast_mean_values.values

    forecast_mean = pd.Series(forecast_mean_values, index=self.test.index)
    self.train.rename('Ret_Train', inplace=True)
    self.test.rename('Ret_Test', inplace=True)
    forecast_mean.rename('Ret_Forecast', inplace=True)
    forecast_vs_test = pd.concat([self.train, self.test, forecast_mean], axis=1)

    forecast_vs_test['Ret_Test_Broad'] = forecast_vs_test[['Ret_Train', 'Ret_Test']].max(axis=1)

    forecast_vs_test['Vol_Train'] = forecast_vs_test[['Ret_Train']].rolling(30).std().shift(1) #30d std dev
    forecast_vs_test['Vol_Test_Broad'] = forecast_vs_test[['Ret_Test_Broad']].rolling(30).std().shift(1) #30d std dev
    forecast_vs_test['Vol_Test'] = np.where(forecast_vs_test['Ret_Test'].isna(), forecast_vs_test['Ret_Test_Broad'], forecast_vs_test['Vol_Test_Broad'])

#     forecast_vol_raw = self.forecast.variance**0.5
#     if self.approach == 'fixed':
#         forecast_std_values = forecast_vol_raw.iloc[0].values
#     elif self.approach == 'roll1d':
#         forecast_std_values = forecast_vol_raw.drop(exclusion_dates)
#         forecast_std_values = forecast_std_values.iloc[:,0].values

    forecast_std_values = self.forecast_std_values.values

    forecast_vs_test_condvol = pd.Series(forecast_std_values, index=self.test.index) #conditional vol
    forecast_vs_test_condvol.rename('Vol_Forecast', inplace=True)
    forecast_vs_test = pd.concat([forecast_vs_test, forecast_vs_test_condvol], axis=1)

    return forecast_vs_test

def forecast_vs_test_mean_plot(self):
    self.forecast_vs_test[['Ret_Train', 'Ret_Test', 'Ret_Forecast']].plot()

def forecast_vs_test_vol_plot(self):
    self.forecast_vs_test[['Vol_Train', 'Vol_Test', 'Vol_Forecast']].plot()

```

```

In [ ]: ar1garch11 = garch('roll1d', 0.9, 1, 1, 'AR', returns, 'Ret_JNJ', 1)
ar1garch12 = garch('roll1d', 0.9, 1, 2, 'AR', returns, 'Ret_JNJ', 1)
ar1garch21 = garch('roll1d', 0.9, 2, 1, 'AR', returns, 'Ret_JNJ', 1)
ar4garch11 = garch('roll1d', 0.9, 1, 1, 'AR', returns, 'Ret_JNJ', 4)
ar4garch12 = garch('roll1d', 0.9, 1, 2, 'AR', returns, 'Ret_JNJ', 4)
ar4garch21 = garch('roll1d', 0.9, 2, 1, 'AR', returns, 'Ret_JNJ', 4)
ar6garch11 = garch('roll1d', 0.9, 1, 1, 'AR', returns, 'Ret_JNJ', 6)
ar6garch12 = garch('roll1d', 0.9, 1, 2, 'AR', returns, 'Ret_JNJ', 6)
ar6garch21 = garch('roll1d', 0.9, 2, 1, 'AR', returns, 'Ret_JNJ', 6)
ar9garch11 = garch('roll1d', 0.9, 1, 1, 'AR', returns, 'Ret_JNJ', 9)
ar9garch12 = garch('roll1d', 0.9, 1, 2, 'AR', returns, 'Ret_JNJ', 9)
ar9garch21 = garch('roll1d', 0.9, 2, 1, 'AR', returns, 'Ret_JNJ', 9)

```

```

Iteration:    1,  Func. Count:    7,  Neg. LLF: 61052.79280867893
Iteration:    2,  Func. Count:   17,  Neg. LLF: 6453202.524870855
Iteration:    3,  Func. Count:   26,  Neg. LLF: 2149143.382759834
Iteration:    4,  Func. Count:   33,  Neg. LLF: 3517.4363591318224
Iteration:    5,  Func. Count:   40,  Neg. LLF: 3726.2849913105483

```



Iteration: 6, Func. Count: 47, Neg. LLF: 3474.4284763938213  
Iteration: 7, Func. Count: 53, Neg. LLF: 3468.286028036266  
Iteration: 8, Func. Count: 59, Neg. LLF: 3466.453192727503  
Iteration: 9, Func. Count: 65, Neg. LLF: 3466.4588223039746  
Iteration: 10, Func. Count: 72, Neg. LLF: 3465.8531836189095  
Iteration: 11, Func. Count: 78, Neg. LLF: 3465.8516812260423  
Iteration: 12, Func. Count: 84, Neg. LLF: 3465.851670324504  
Iteration: 13, Func. Count: 90, Neg. LLF: 3465.8516689144944  
Iteration: 14, Func. Count: 95, Neg. LLF: 3465.8516682923864

Optimization terminated successfully (Exit mode 0)  
Current function value: 3465.8516689144944  
Iterations: 14

Function evaluations: 95  
Gradient evaluations: 14

Iteration: 1, Func. Count: 8, Neg. LLF: 69083.69358667244  
Iteration: 2, Func. Count: 19, Neg. LLF: 653012.6043595799  
Iteration: 3, Func. Count: 29, Neg. LLF: 194456.95020774924  
Iteration: 4, Func. Count: 37, Neg. LLF: 3501.82492811266  
Iteration: 5, Func. Count: 44, Neg. LLF: 4090.229723472999  
Iteration: 6, Func. Count: 55, Neg. LLF: 5942.46722337079  
Iteration: 7, Func. Count: 63, Neg. LLF: 3509.8033411977813  
Iteration: 8, Func. Count: 71, Neg. LLF: 6261.042107955817  
Iteration: 9, Func. Count: 79, Neg. LLF: 3539.4626232689297  
Iteration: 10, Func. Count: 87, Neg. LLF: 3467.9994882246  
Iteration: 11, Func. Count: 94, Neg. LLF: 3465.858098309489  
Iteration: 12, Func. Count: 101, Neg. LLF: 3464.918364053072  
Iteration: 13, Func. Count: 108, Neg. LLF: 3464.641273723206  
Iteration: 14, Func. Count: 115, Neg. LLF: 3464.271482345731  
Iteration: 15, Func. Count: 122, Neg. LLF: 3469.1688783718078  
Iteration: 16, Func. Count: 130, Neg. LLF: 3468.759623816978  
Iteration: 17, Func. Count: 138, Neg. LLF: 3464.1262951008443  
Iteration: 18, Func. Count: 146, Neg. LLF: 3463.3881920539093  
Iteration: 19, Func. Count: 153, Neg. LLF: 3463.3784688239493  
Iteration: 20, Func. Count: 160, Neg. LLF: 3463.3778556887087  
Iteration: 21, Func. Count: 167, Neg. LLF: 3463.377790240138  
Iteration: 22, Func. Count: 174, Neg. LLF: 3463.377759874371  
Iteration: 23, Func. Count: 181, Neg. LLF: 3463.377758124896  
Iteration: 24, Func. Count: 187, Neg. LLF: 3463.377757729458

Optimization terminated successfully (Exit mode 0)  
Current function value: 3463.377758124896  
Iterations: 24

Function evaluations: 187  
Gradient evaluations: 24

Iteration: 1, Func. Count: 8, Neg. LLF: 17937.57810219515  
Iteration: 2, Func. Count: 19, Neg. LLF: 12067.049581800402  
Iteration: 3, Func. Count: 29, Neg. LLF: 4308684.068809033  
Iteration: 4, Func. Count: 37, Neg. LLF: 4324.496639023809  
Iteration: 5, Func. Count: 45, Neg. LLF: 4085.5710164266698  
Iteration: 6, Func. Count: 53, Neg. LLF: 3477.0671584381316  
Iteration: 7, Func. Count: 60, Neg. LLF: 3557.145130122208  
Iteration: 8, Func. Count: 68, Neg. LLF: 3468.412369435724  
Iteration: 9, Func. Count: 75, Neg. LLF: 3466.3265758692205  
Iteration: 10, Func. Count: 82, Neg. LLF: 3465.89949425693  
Iteration: 11, Func. Count: 89, Neg. LLF: 3465.859659552012  
Iteration: 12, Func. Count: 96, Neg. LLF: 3465.8520416145147  
Iteration: 13, Func. Count: 103, Neg. LLF: 3465.851669786107  
Iteration: 14, Func. Count: 110, Neg. LLF: 3465.8516689141215

Optimization terminated successfully (Exit mode 0)  
Current function value: 3465.8516689141215  
Iterations: 14

Function evaluations: 110  
Gradient evaluations: 14

Iteration: 1, Func. Count: 10, Neg. LLF: 107330.04160288133  
Iteration: 2, Func. Count: 23, Neg. LLF: 266139.92923819047  
Iteration: 3, Func. Count: 35, Neg. LLF: 2989524.886098926  
Iteration: 4, Func. Count: 47, Neg. LLF: 3318250.735340111  
Iteration: 5, Func. Count: 57, Neg. LLF: 536188.7902039221  
Iteration: 6, Func. Count: 67, Neg. LLF: 9356.110124948276  
Iteration: 7, Func. Count: 78, Neg. LLF: 33807.1634548685  
Iteration: 8, Func. Count: 90, Neg. LLF: 3665.12416696479  
Iteration: 9, Func. Count: 100, Neg. LLF: 3644.5712264644308  
Iteration: 10, Func. Count: 110, Neg. LLF: 3449.0239295454876  
Iteration: 11, Func. Count: 119, Neg. LLF: 3438.380658383642  
Iteration: 12, Func. Count: 128, Neg. LLF: 3437.5974908116596  
Iteration: 13, Func. Count: 137, Neg. LLF: 3437.5499037554855  
Iteration: 14, Func. Count: 146, Neg. LLF: 3437.540667237796  
Iteration: 15, Func. Count: 155, Neg. LLF: 3437.539538200621  
Iteration: 16, Func. Count: 164, Neg. LLF: 3437.539503588682  
Iteration: 17, Func. Count: 173, Neg. LLF: 3437.5395029441615

Optimization terminated successfully (Exit mode 0)  
Current function value: 3437.5395029441615  
Iterations: 17

Function evaluations: 173

Gradient evaluations: 17  
Iteration: 1, Func. Count: 11, Neg. LLF: 105720.44470005757  
Iteration: 2, Func. Count: 25, Neg. LLF: 683766.6679308289  
Iteration: 3, Func. Count: 38, Neg. LLF: 5462395.862088518  
Iteration: 4, Func. Count: 51, Neg. LLF: 1444737.1997767761  
Iteration: 5, Func. Count: 62, Neg. LLF: 638258.2652303372  
Iteration: 6, Func. Count: 73, Neg. LLF: 29198.0853324265  
Iteration: 7, Func. Count: 86, Neg. LLF: 22402.62799317786  
Iteration: 8, Func. Count: 99, Neg. LLF: 3623.1679622507577  
Iteration: 9, Func. Count: 110, Neg. LLF: 3714.667514908565  
Iteration: 10, Func. Count: 121, Neg. LLF: 3526.277068510915  
Iteration: 11, Func. Count: 132, Neg. LLF: 3433.865890560791  
Iteration: 12, Func. Count: 142, Neg. LLF: 3512.7631667158885  
Iteration: 13, Func. Count: 153, Neg. LLF: 7170.226133384282  
Iteration: 14, Func. Count: 164, Neg. LLF: 3427.9876264156064  
Iteration: 15, Func. Count: 174, Neg. LLF: 3427.692363021994  
Iteration: 16, Func. Count: 184, Neg. LLF: 3427.681634449081  
Iteration: 17, Func. Count: 194, Neg. LLF: 3427.6800815500483  
Iteration: 18, Func. Count: 204, Neg. LLF: 3427.6788795949988  
Iteration: 19, Func. Count: 214, Neg. LLF: 3427.6787273519885  
Iteration: 20, Func. Count: 224, Neg. LLF: 3427.678719087139  
Iteration: 21, Func. Count: 233, Neg. LLF: 3427.6787182690528

Optimization terminated successfully (Exit mode 0)

Current function value: 3427.678719087139

Iterations: 21

Function evaluations: 233

Gradient evaluations: 21

Iteration: 1, Func. Count: 11, Neg. LLF: 108299.77978454804  
Iteration: 2, Func. Count: 25, Neg. LLF: 19340.934776461952  
Iteration: 3, Func. Count: 38, Neg. LLF: 1210590.1341135735  
Iteration: 4, Func. Count: 50, Neg. LLF: 1782770.466980549  
Iteration: 5, Func. Count: 61, Neg. LLF: 7005.79502068715  
Iteration: 6, Func. Count: 73, Neg. LLF: 10650.717901502567  
Iteration: 7, Func. Count: 85, Neg. LLF: 14875.376959228373  
Iteration: 8, Func. Count: 98, Neg. LLF: 4056.0141719267535  
Iteration: 9, Func. Count: 109, Neg. LLF: 3461.7022728077327  
Iteration: 10, Func. Count: 119, Neg. LLF: 3509.3295800203277  
Iteration: 11, Func. Count: 130, Neg. LLF: 3459.7770224371625  
Iteration: 12, Func. Count: 141, Neg. LLF: 3438.0374668686536  
Iteration: 13, Func. Count: 151, Neg. LLF: 3438.575662515697  
Iteration: 14, Func. Count: 162, Neg. LLF: 3437.592100089162  
Iteration: 15, Func. Count: 172, Neg. LLF: 3437.547091149  
Iteration: 16, Func. Count: 182, Neg. LLF: 3437.539624467403  
Iteration: 17, Func. Count: 192, Neg. LLF: 3437.5395061532554  
Iteration: 18, Func. Count: 202, Neg. LLF: 3437.5395029536576  
Iteration: 19, Func. Count: 211, Neg. LLF: 3437.5395020398882

Optimization terminated successfully (Exit mode 0)

Current function value: 3437.5395029536576

Iterations: 19

Function evaluations: 211

Gradient evaluations: 19

Iteration: 1, Func. Count: 12, Neg. LLF: 86800.32633580668  
Iteration: 2, Func. Count: 27, Neg. LLF: 345174.4364261344  
Iteration: 3, Func. Count: 41, Neg. LLF: 1823096.5120845647  
Iteration: 4, Func. Count: 55, Neg. LLF: 991167.6355774499  
Iteration: 5, Func. Count: 67, Neg. LLF: 2476266.307154471  
Iteration: 6, Func. Count: 79, Neg. LLF: 23587.09144941823  
Iteration: 7, Func. Count: 93, Neg. LLF: 524973.7651144749  
Iteration: 8, Func. Count: 107, Neg. LLF: 258862.99512491532  
Iteration: 9, Func. Count: 120, Neg. LLF: 654049.8307928008  
Iteration: 10, Func. Count: 134, Neg. LLF: 8823.155759110101  
Iteration: 11, Func. Count: 148, Neg. LLF: 3523.9009235776066  
Iteration: 12, Func. Count: 160, Neg. LLF: 3460.7104580802793  
Iteration: 13, Func. Count: 172, Neg. LLF: 3432.237765958468  
Iteration: 14, Func. Count: 183, Neg. LLF: 3430.693247739154  
Iteration: 15, Func. Count: 194, Neg. LLF: 3430.624392349625  
Iteration: 16, Func. Count: 205, Neg. LLF: 3430.617922502993  
Iteration: 17, Func. Count: 216, Neg. LLF: 3430.6175694173253  
Iteration: 18, Func. Count: 227, Neg. LLF: 3430.6175498167263  
Iteration: 19, Func. Count: 238, Neg. LLF: 3430.617547419066  
Iteration: 20, Func. Count: 248, Neg. LLF: 3430.6175465758624

Optimization terminated successfully (Exit mode 0)

Current function value: 3430.617547419066

Iterations: 20

Function evaluations: 248

Gradient evaluations: 20

Iteration: 1, Func. Count: 13, Neg. LLF: 594183.6739543932  
Iteration: 2, Func. Count: 29, Neg. LLF: 548510.0915971808  
Iteration: 3, Func. Count: 44, Neg. LLF: 652507.7977986201  
Iteration: 4, Func. Count: 59, Neg. LLF: 1014958.531062253  
Iteration: 5, Func. Count: 72, Neg. LLF: 1333871.4017855674  
Iteration: 6, Func. Count: 85, Neg. LLF: 16055.98701418603  
Iteration: 7, Func. Count: 99, Neg. LLF: 14393.760649319356

Iteration: 8, Func. Count: 114, Neg. LLF: 32802.56181942888  
Iteration: 9, Func. Count: 129, Neg. LLF: 6811.624067239947  
Iteration: 10, Func. Count: 143, Neg. LLF: 3672.198131222566  
Iteration: 11, Func. Count: 156, Neg. LLF: 4748.203685468032  
Iteration: 12, Func. Count: 170, Neg. LLF: 3438.05532187779  
Iteration: 13, Func. Count: 182, Neg. LLF: 3439.04765217898  
Iteration: 14, Func. Count: 196, Neg. LLF: 7731.157719820667  
Iteration: 15, Func. Count: 209, Neg. LLF: 3452.7187814758854  
Iteration: 16, Func. Count: 222, Neg. LLF: 3417.864598077047  
Iteration: 17, Func. Count: 234, Neg. LLF: 3417.4373825750017  
Iteration: 18, Func. Count: 246, Neg. LLF: 3417.264599118532  
Iteration: 19, Func. Count: 258, Neg. LLF: 3417.2367763351685  
Iteration: 20, Func. Count: 270, Neg. LLF: 3417.2340839065464  
Iteration: 21, Func. Count: 282, Neg. LLF: 3417.234044639953  
Iteration: 22, Func. Count: 294, Neg. LLF: 3417.2340426349983  
Iteration: 23, Func. Count: 305, Neg. LLF: 3417.234041627248

Optimization terminated successfully (Exit mode 0)

Current function value: 3417.2340426349983

Iterations: 23

Function evaluations: 305

Gradient evaluations: 23

Iteration: 1, Func. Count: 13, Neg. LLF: 75218.41200646257  
Iteration: 2, Func. Count: 29, Neg. LLF: 19671.872310762952  
Iteration: 3, Func. Count: 44, Neg. LLF: 2268448.796854481  
Iteration: 4, Func. Count: 58, Neg. LLF: 3291596.9174673483  
Iteration: 5, Func. Count: 71, Neg. LLF: 7430.407614743193  
Iteration: 6, Func. Count: 85, Neg. LLF: 13140.885411368079  
Iteration: 7, Func. Count: 99, Neg. LLF: 16859.28419801085  
Iteration: 8, Func. Count: 114, Neg. LLF: 9295.661133562644  
Iteration: 9, Func. Count: 129, Neg. LLF: 5161.525684388116  
Iteration: 10, Func. Count: 142, Neg. LLF: 4590.076086454085  
Iteration: 11, Func. Count: 155, Neg. LLF: 3702.851821141414  
Iteration: 12, Func. Count: 168, Neg. LLF: 3512.3050201709857  
Iteration: 13, Func. Count: 181, Neg. LLF: 3433.4192110069507  
Iteration: 14, Func. Count: 193, Neg. LLF: 3433.777360688722  
Iteration: 15, Func. Count: 206, Neg. LLF: 3430.697009006723  
Iteration: 16, Func. Count: 218, Neg. LLF: 3430.6289925951  
Iteration: 17, Func. Count: 230, Neg. LLF: 3430.619013641425  
Iteration: 18, Func. Count: 242, Neg. LLF: 3430.6176832330766  
Iteration: 19, Func. Count: 254, Neg. LLF: 3430.6175526901698  
Iteration: 20, Func. Count: 266, Neg. LLF: 3430.6175476085923  
Iteration: 21, Func. Count: 277, Neg. LLF: 3430.617546765311

Optimization terminated successfully (Exit mode 0)

Current function value: 3430.6175476085923

Iterations: 21

Function evaluations: 277

Gradient evaluations: 21

Iteration: 1, Func. Count: 15, Neg. LLF: 88707.84092169083  
Iteration: 2, Func. Count: 34, Neg. LLF: 24352265.54097928  
Iteration: 3, Func. Count: 51, Neg. LLF: 7423026.486659714  
Iteration: 4, Func. Count: 68, Neg. LLF: 2132723.1804170315  
Iteration: 5, Func. Count: 84, Neg. LLF: 1128168.5976529364  
Iteration: 6, Func. Count: 99, Neg. LLF: 11516930.5615177  
Iteration: 7, Func. Count: 115, Neg. LLF: 5369893.44731659  
Iteration: 8, Func. Count: 132, Neg. LLF: 24503.60825122107  
Iteration: 9, Func. Count: 149, Neg. LLF: 17890.22923049912  
Iteration: 10, Func. Count: 166, Neg. LLF: 24604.389959691376  
Iteration: 11, Func. Count: 183, Neg. LLF: 5591.345220345723  
Iteration: 12, Func. Count: 199, Neg. LLF: 5435.762395135499  
Iteration: 13, Func. Count: 215, Neg. LLF: 15794.771730555112  
Iteration: 14, Func. Count: 232, Neg. LLF: 3713.6022761414097  
Iteration: 15, Func. Count: 247, Neg. LLF: 4216.03661643017  
Iteration: 16, Func. Count: 262, Neg. LLF: 3409.5223577266306  
Iteration: 17, Func. Count: 276, Neg. LLF: 3407.074053427291  
Iteration: 18, Func. Count: 290, Neg. LLF: 3405.1217499163035  
Iteration: 19, Func. Count: 304, Neg. LLF: 3405.018856448305  
Iteration: 20, Func. Count: 318, Neg. LLF: 3405.015295212379  
Iteration: 21, Func. Count: 332, Neg. LLF: 3405.014895329883  
Iteration: 22, Func. Count: 346, Neg. LLF: 3405.0148767988976  
Iteration: 23, Func. Count: 360, Neg. LLF: 3405.014874665576  
Iteration: 24, Func. Count: 373, Neg. LLF: 3405.014873426846

Optimization terminated successfully (Exit mode 0)

Current function value: 3405.014874665576

Iterations: 24

Function evaluations: 373

Gradient evaluations: 24

Iteration: 1, Func. Count: 16, Neg. LLF: 142408.81828828537  
Iteration: 2, Func. Count: 36, Neg. LLF: 11821321.537696928  
Iteration: 3, Func. Count: 54, Neg. LLF: 141619741.40669888  
Iteration: 4, Func. Count: 72, Neg. LLF: 1199931.9879678087  
Iteration: 5, Func. Count: 89, Neg. LLF: 2405569.4319773437  
Iteration: 6, Func. Count: 105, Neg. LLF: 17263027.95102325  
Iteration: 7, Func. Count: 122, Neg. LLF: 17937.852518408086

```

Iteration:      8,  Func. Count:    140,  Neg. LLF: 13584.148050944752
Iteration:      9,  Func. Count:    158,  Neg. LLF: 17029.57502359087
Iteration:     10,  Func. Count:    176,  Neg. LLF: 8388.460968971425
Iteration:     11,  Func. Count:    193,  Neg. LLF: 3867.7838429741555
Iteration:     12,  Func. Count:    209,  Neg. LLF: 6298.533612131359
Iteration:     13,  Func. Count:    226,  Neg. LLF: 11531.41715935291
Iteration:     14,  Func. Count:    244,  Neg. LLF: 4900353975.574669
Iteration:     15,  Func. Count:    261,  Neg. LLF: 3622.7469841542415
Iteration:     16,  Func. Count:    277,  Neg. LLF: 3405.480493787151
Iteration:     17,  Func. Count:    292,  Neg. LLF: 3433.6159799671827
Iteration:     18,  Func. Count:    308,  Neg. LLF: 4086.271309528178
Iteration:     19,  Func. Count:    325,  Neg. LLF: 3416.3648966115884
Iteration:     20,  Func. Count:    341,  Neg. LLF: 3395.121305412791
Iteration:     21,  Func. Count:    356,  Neg. LLF: 3395.0766436693
Iteration:     22,  Func. Count:    371,  Neg. LLF: 3395.064650251343
Iteration:     23,  Func. Count:    386,  Neg. LLF: 3395.0580011683596
Iteration:     24,  Func. Count:    401,  Neg. LLF: 3395.0577128220684
Iteration:     25,  Func. Count:    416,  Neg. LLF: 3395.057682108208
Iteration:     26,  Func. Count:    431,  Neg. LLF: 3395.05768077473
Iteration:     27,  Func. Count:    445,  Neg. LLF: 3395.0576794840154

```

Optimization terminated successfully (Exit mode 0)

Current function value: 3395.05768077473

Iterations: 27

Function evaluations: 445

Gradient evaluations: 27

```

Iteration:      1,  Func. Count:     16,  Neg. LLF: 105829.2646551054
Iteration:      2,  Func. Count:     35,  Neg. LLF: 4254697.025616312
Iteration:      3,  Func. Count:     53,  Neg. LLF: 37468110.06297142
Iteration:      4,  Func. Count:     70,  Neg. LLF: 2235955.244480608
Iteration:      5,  Func. Count:     86,  Neg. LLF: 2801106.530496165
Iteration:      6,  Func. Count:    102,  Neg. LLF: 40680.993817821
Iteration:      7,  Func. Count:    120,  Neg. LLF: 22003.56627718916
Iteration:      8,  Func. Count:    138,  Neg. LLF: 126849658.46430635
Iteration:      9,  Func. Count:    156,  Neg. LLF: 3036136.989924136
Iteration:     10,  Func. Count:    173,  Neg. LLF: 53284242.370206356
Iteration:     11,  Func. Count:    190,  Neg. LLF: 15391.939605237068
Iteration:     12,  Func. Count:    208,  Neg. LLF: 17067.173799006396
Iteration:     13,  Func. Count:    226,  Neg. LLF: 6199.368876048974
Iteration:     14,  Func. Count:    243,  Neg. LLF: 3908.2730870805417
Iteration:     15,  Func. Count:    259,  Neg. LLF: 4032.453741903061
Iteration:     16,  Func. Count:    275,  Neg. LLF: 3414.4177276076225
Iteration:     17,  Func. Count:    290,  Neg. LLF: 3405.685118540223
Iteration:     18,  Func. Count:    305,  Neg. LLF: 3405.110160567734
Iteration:     19,  Func. Count:    320,  Neg. LLF: 3405.0275977411047
Iteration:     20,  Func. Count:    335,  Neg. LLF: 3405.0156115273357
Iteration:     21,  Func. Count:    350,  Neg. LLF: 3405.0149199813623
Iteration:     22,  Func. Count:    365,  Neg. LLF: 3405.014876521888
Iteration:     23,  Func. Count:    380,  Neg. LLF: 3405.0148746674645
Iteration:     24,  Func. Count:    394,  Neg. LLF: 3405.0148734289187

```

Optimization terminated successfully (Exit mode 0)

Current function value: 3405.0148746674645

Iterations: 24

Function evaluations: 394

Gradient evaluations: 24

```

In [ ]: runs = [ar1garch11, ar1garch12, ar1garch21, ar4garch11, ar4garch12, ar4garch21, ar6garch11, ar6garch12, ar6garch21]

for i, ic in enumerate(runs):
    ic_append = ic.info_crit()
    ic_append['Model'] = ic.specs()

    if i == 0:
        ics = ic_append
    else:
        ics = pd.concat([ics, ic_append])

cols = ics.columns.tolist()
cols = cols[:-1] + cols[-1:]
ics = ics[cols]
ics

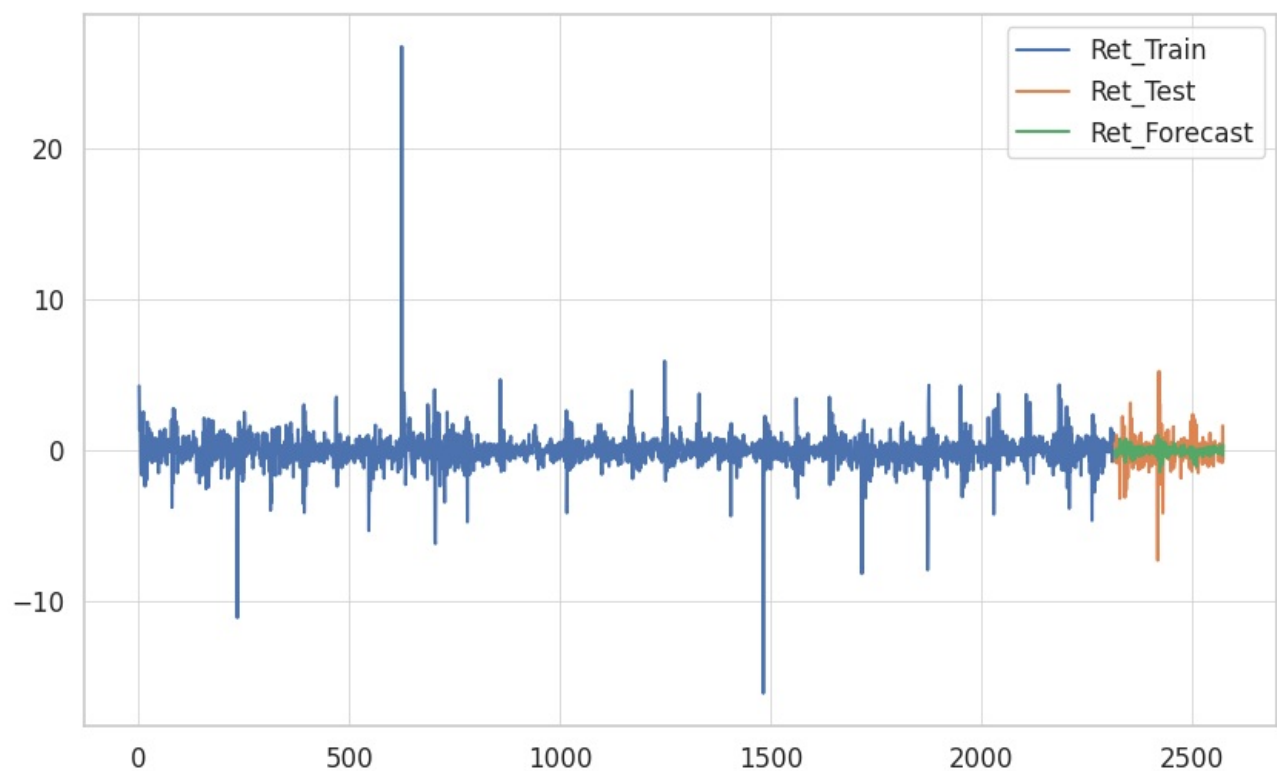
```

Out[ ]:	Model	AIC	BIC
0	ar1garch11	6941.703338	6970.437002
0	ar1garch12	6938.755516	6973.235914
0	ar1garch21	6943.703338	6978.183735
0	ar4garch11	6891.079006	6937.042491
0	ar4garch12	6873.357438	6925.066359
0	ar4garch21	6893.079006	6944.787926
0	ar6garch11	6881.235095	6938.680793
0	ar6garch12	6856.468085	6919.658353
0	ar6garch21	6883.235095	6946.425363
0	ar9garch11	6836.029749	6910.692255
0	ar9garch12	6818.115362	6898.521137
0	ar9garch21	6838.029749	6918.435525

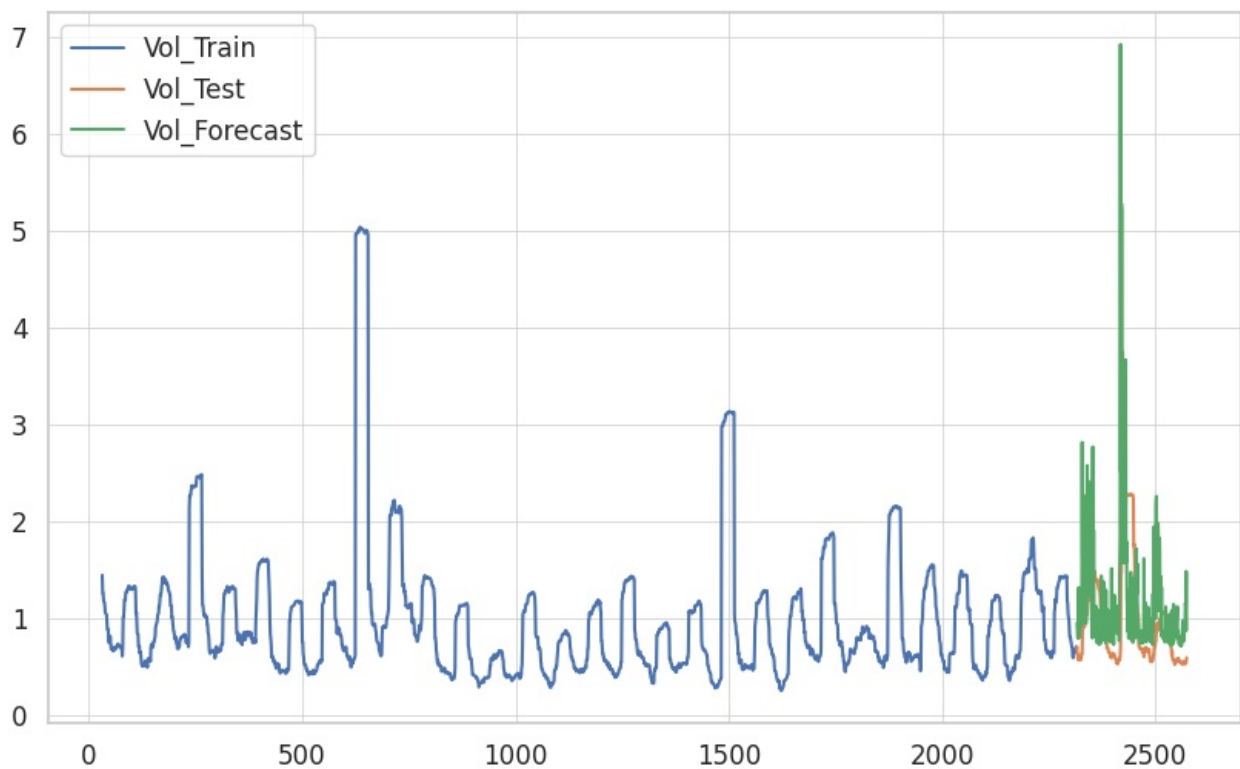
It can be seen that across all model variations AR(9)-GARCH(1,2) had the lowest Akaike Information Criterion (AIC) and was the model of choice.

Representations of its mean and volatility forecasts are shown below.

```
In [ ]: ar9garch12.forecast_vs_test_mean_plot()
```



```
In [ ]: ar9garch12.forecast_vs_test_vol_plot()
```



## Strategy & Results

The forecasts were used to generate trading signals. Whenever the next period return was forecast to be positive, a long position was selected, otherwise a short position was selected. The position was only initiated if the next period volatility was forecast to be above the last observed volatility level, i.e. if a relatively large move in either direction was anticipated to increase the chances of profiting from the trade and avoiding unnecessary transaction costs.

Gross P&L was calculated based on the direction and realised return. Net P&L subtracted trading costs (5 basis points of trade value) from this wherever there was a change in position day-on-day.

```
In [ ]: class strategy(garch):
    def __init__(self, proportion, approach, p, q, mean, data, asset, lags, o = 0):
        garch.__init__(self, proportion, approach, p, q, mean, data, asset, lags, o = 0)
        self.trade = self.build()

    def build(self):
        trade = self.forecast_vs_test
        print(self.last_train_date)

        close = self.data[['JNJ']]
        trade = pd.merge(trade, close, left_index = True, right_index = True)

        last_train = trade.iloc[lambdax: x.index == self.last_train_date]
        last_vol_train = last_train['Vol_Train'].iloc[0]

        trade['direction'] = np.where(trade['Ret_Train'].notna(), 0, np.where(trade['Ret_Forecast'] > 0, 1, -1))
        trade['amplifier'] = np.where(trade['Ret_Train'].notna(), 0, np.where(trade['Vol_Forecast'] > last_vol_train, 1, -1))
        trade['position'] = trade['direction'] * trade['amplifier']
        trade['pnl_gross'] = trade['Ret_Test'] / 1000 * trade['position'] # adjust returns back to normal level
        trade['pnl_gross_cum'] = trade['pnl_gross'].cumsum()
        trade['position_lag'] = trade['position'].shift(1)
        trade['cost'] = np.where(trade['position'] == trade['position_lag'], 0, TRANSACTION_COSTS * trade['JNJ'])
        trade['pnl_net'] = trade['pnl_gross'] - trade['cost']
        trade['pnl_net_cum'] = trade['pnl_net'].cumsum()

        trade = trade.iloc[lambdax: x.index >= self.last_train_date]

        return trade

    def pos(self):
        plt.figure(figsize=(16, 8), dpi=150)

        self.trade['direction'].plot(label='Direction')
        self.trade['amplifier'].plot(label='Amplifier')
        self.trade['position'].plot(label='Position')
        self.trade['cost'].plot(label='Cost')

        plt.title('Trading Positions ' + self.asset + ' ' + self.specs())
        plt.legend()
```

```

def pnl_gross(self):
    plt.figure(figsize=(16, 8), dpi=150)

    self.trade['pnl_gross'].plot(label='Gross P&L')
    self.trade['pnl_gross_cum'].plot(label='Cumulative Gross P&L')

    plt.title('Trading Gross P&L ' + self.asset + ' ' + self.specs())
    plt.legend()

def pnl_net(self):
    plt.figure(figsize=(16, 8), dpi=150)

    self.trade['pnl_net'].plot(label='Net P&L')
    self.trade['pnl_net_cum'].plot(label='Cumulative Net P&L')

    plt.title('Trading Net P&L ' + self.asset + ' ' + self.specs())
    plt.legend()

def sr(self):
    sr_gross = 252 ** 0.5 * np.mean(self.trade['pnl_gross']) / np.std(self.trade['pnl_gross'])
    sr_net = 252 ** 0.5 * np.mean(self.trade['pnl_net']) / np.std(self.trade['pnl_net'])

    return pd.DataFrame({'Metric': ['SR Gross', 'SR Net'], 'Value': [sr_gross, sr_net]})

```

```
In [ ]: ar9garch12_strat = strategy('roll1d', 0.9, 1, 2, 'AR', returns, 'Ret_JNJ', 9)
```

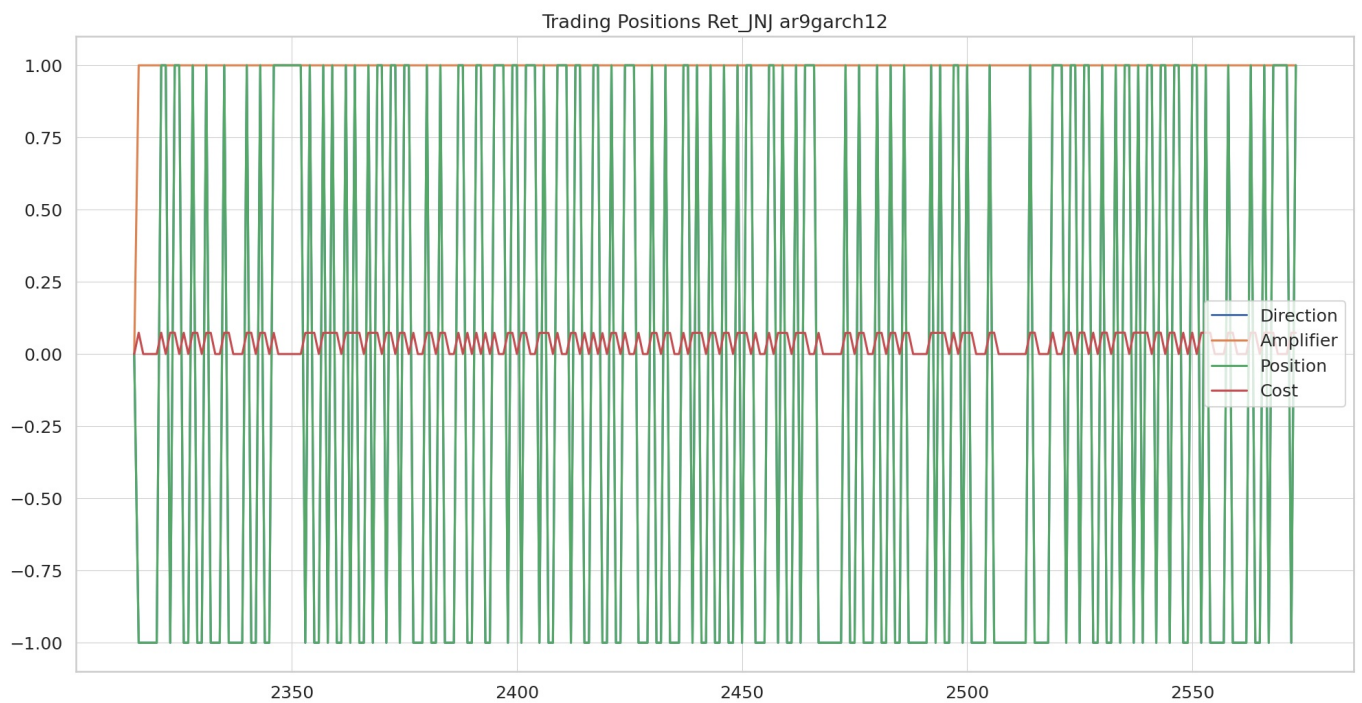
```

Iteration:      1,  Func. Count:      16,  Neg. LLF: 142408.81828828537
Iteration:      2,  Func. Count:      36,  Neg. LLF: 11821321.537696928
Iteration:      3,  Func. Count:      54,  Neg. LLF: 141619741.40669888
Iteration:      4,  Func. Count:      72,  Neg. LLF: 1199931.9879678087
Iteration:      5,  Func. Count:      89,  Neg. LLF: 2405569.4319773437
Iteration:      6,  Func. Count:     105,  Neg. LLF: 17263027.95102325
Iteration:      7,  Func. Count:     122,  Neg. LLF: 17937.852518408086
Iteration:      8,  Func. Count:     140,  Neg. LLF: 13584.148050944752
Iteration:      9,  Func. Count:     158,  Neg. LLF: 17029.57502359087
Iteration:     10,  Func. Count:     176,  Neg. LLF: 8388.460968971425
Iteration:     11,  Func. Count:     193,  Neg. LLF: 3867.7838429741555
Iteration:     12,  Func. Count:     209,  Neg. LLF: 6298.533612131359
Iteration:     13,  Func. Count:     226,  Neg. LLF: 11531.41715935291
Iteration:     14,  Func. Count:     244,  Neg. LLF: 4900353975.574669
Iteration:     15,  Func. Count:     261,  Neg. LLF: 3622.7469841542415
Iteration:     16,  Func. Count:     277,  Neg. LLF: 3405.480493787151
Iteration:     17,  Func. Count:     292,  Neg. LLF: 3433.6159799671827
Iteration:     18,  Func. Count:     308,  Neg. LLF: 4086.271309528178
Iteration:     19,  Func. Count:     325,  Neg. LLF: 3416.3648966115884
Iteration:     20,  Func. Count:     341,  Neg. LLF: 3395.121305412791
Iteration:     21,  Func. Count:     356,  Neg. LLF: 3395.0766436693
Iteration:     22,  Func. Count:     371,  Neg. LLF: 3395.064650251343
Iteration:     23,  Func. Count:     386,  Neg. LLF: 3395.0580011683596
Iteration:     24,  Func. Count:     401,  Neg. LLF: 3395.0577128220684
Iteration:     25,  Func. Count:     416,  Neg. LLF: 3395.057682108208
Iteration:     26,  Func. Count:     431,  Neg. LLF: 3395.05768077473
Iteration:     27,  Func. Count:     445,  Neg. LLF: 3395.0576794840154
Optimization terminated successfully (Exit mode 0)
    Current function value: 3395.05768077473
    Iterations: 27
    Function evaluations: 445
    Gradient evaluations: 27

```

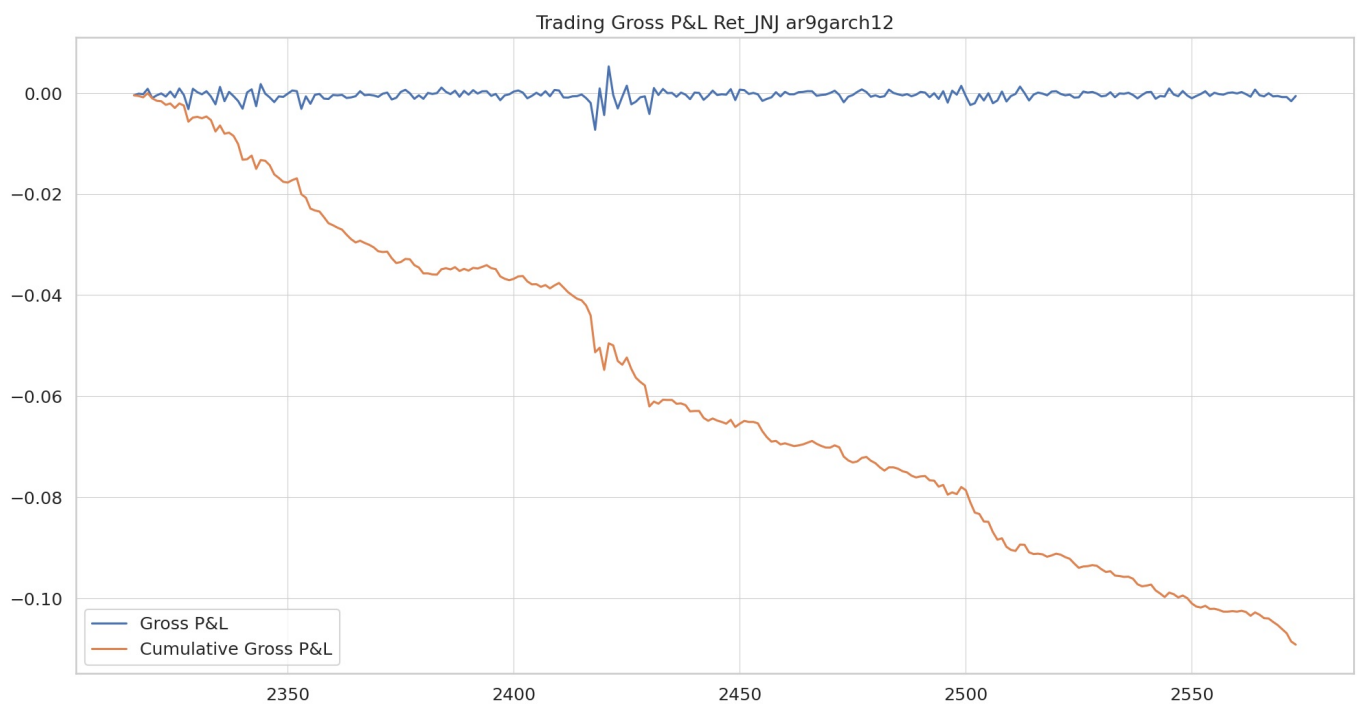
2315

```
In [ ]: ar9garch12_strat.pos()
```



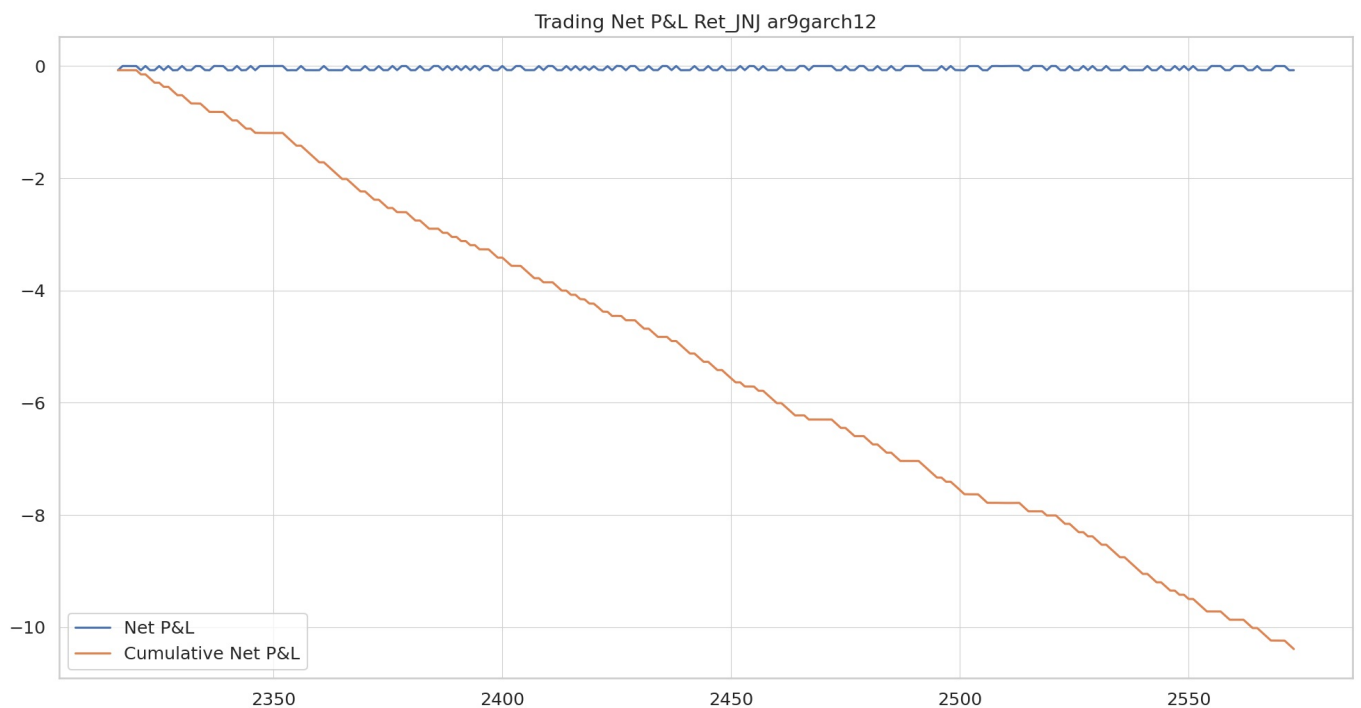
The strategy resulted in frequent trades, often switching positions between long and short. These were always triggered due to the volatility forecast always exceeding the last observed level of volatility. The volatility forecast was inaccurate in line with lack of arch effects and no statistically significant squared log returns observed earlier. The frequent execution of trades implied high trading costs.

```
In [ ]: ar9garch12_strat.pnl_gross()
```



```
In [ ]: ar9garch12_strat.pnl_net()
```





As expected, net cumulative P&L was negative. Gross cumulative P&L was also negative, which suggests being in the market (momentum strategy) based on recently observed returns is not a profitable strategy. A mean-reversion strategy may have performed better.

```
In [ ]: ar9garch12_strat.sr()
```

```
Out[ ]:
```

	Metric	Value
0	SR Gross	-6.581488
1	SR Net	-17.505621

As a result, the net Sharpe ratio on the test set was -17.5.

## Data: LSTM

LSTM (Long Short Term Memory) was the second model for performing univariate prediction of the time series of Johnson & Johnson equity price.

```
In [ ]: # Shared processing imports start here...
import numpy as np          # General purpose math
import pandas as pd         # Dataframes
import matplotlib.pyplot as plt # Simple plotting
# End of shared processing imports.

# LSTM processing imports start here...
import tensorflow as tf     # LSTM support
from tensorflow.keras.layers import LSTM # LSTM model definition
from tensorflow.keras.models import Sequential # Sequential LSTM layer
from tensorflow.keras.layers import Dense # Dense LSTM layer
from sklearn.preprocessing import MinMaxScaler # Transforming the observations to a specific range
from sklearn.model_selection import train_test_split # Enabling out-of-sample validation
from sklearn.metrics import mean_squared_error # Model evaluation
import quantstats as qs    # calculating ratios
# End of LSTM processing imports.

In [ ]: jnj_close = pd.read_csv('../jnj.us.txt')[["<DATE>", "<TIME>", "<OPEN>", "<CLOSE>"]].rename(columns={'<DATE>': '<TIME>': '<OPEN>': '<CLOSE>':
jnj_close.head()
```

Out [ ]:

	Date	Time	Open	Close
0	20240419	153000	146.150	144.8000
1	20240419	153500	144.800	145.4200
2	20240419	154000	145.440	145.6200
3	20240419	154500	145.660	145.9858
4	20240419	155000	145.975	146.1800

The date and time information was split between two columns. A single column unifying the two was created to allow for later construction of a timestamp.

```
In [ ]: jnj_close['Datetime String'] = jnj_close.Date.map(str) + " " + jnj_close.Time.map(str)
jnj_close.head()
```

Out [ ]:

	Date	Time	Open	Close	Datetime String
0	20240419	153000	146.150	144.8000	20240419 153000
1	20240419	153500	144.800	145.4200	20240419 153500
2	20240419	154000	145.440	145.6200	20240419 154000
3	20240419	154500	145.660	145.9858	20240419 154500
4	20240419	155000	145.975	146.1800	20240419 155000

```
In [ ]: jnj_close['Datetime'] = pd.to_datetime(jnj_close['Datetime String'], format='%Y%m%d %H%M%S')
jnj_close = jnj_close.set_index('Datetime').drop(columns=['Datetime String', 'Date', 'Time'])
jnj_close.head()
```

Out [ ]:

	Open	Close
Datetime		
2024-04-19 15:30:00	146.150	144.8000
2024-04-19 15:35:00	144.800	145.4200
2024-04-19 15:40:00	145.440	145.6200
2024-04-19 15:45:00	145.660	145.9858
2024-04-19 15:50:00	145.975	146.1800

The 'Close' variable was the target for predicting. The prediction model had to operate on lagged open/close variables to ensure look-ahead bias was eliminated.

```
In [ ]: jnj_close['Open Lag'] = jnj_close['Open'].shift(6)
jnj_close['Close Lag'] = jnj_close['Close'].shift(6)
jnj_close.head(10)
```

Out [ ]:

	Open	Close	Open Lag	Close Lag
Datetime				
2024-04-19 15:30:00	146.150	144.8000	NaN	NaN
2024-04-19 15:35:00	144.800	145.4200	NaN	NaN
2024-04-19 15:40:00	145.440	145.6200	NaN	NaN
2024-04-19 15:45:00	145.660	145.9858	NaN	NaN
2024-04-19 15:50:00	145.975	146.1800	NaN	NaN
2024-04-19 15:55:00	146.180	146.0800	NaN	NaN
2024-04-19 16:00:00	146.090	145.8450	146.15	144.8000
2024-04-19 16:05:00	145.830	146.1700	144.80	145.4200
2024-04-19 16:10:00	146.170	146.3000	145.44	145.6200
2024-04-19 16:15:00	146.320	146.1800	145.66	145.9858

```
In [ ]: jnj_close = jnj_close.dropna()
```

# Methodology: LSTM

# Model

The model building process was started by affixing the random seed for reproducibility.

```
In [ ]: tf.random.set_seed(7)
```

LSTMs are sensitive to input data scale. Dataframe values were normalised before fitting the model to ensure best performance.

```
In [ ]: scaler = MinMaxScaler(feature_range=(0, 1))
jnj_close = pd.DataFrame(scaler.fit_transform(jnj_close), index=jnj_close.index, columns=jnj_close.columns)
jnj_close.head()
```

```
Out[ ]:
```

	Open	Close	Open Lag	Close Lag
Datetime				
2024-04-19 16:00:00	0.209603	0.169242	0.215143	0.070714
2024-04-19 16:05:00	0.185596	0.199885	0.090489	0.129171
2024-04-19 16:10:00	0.216990	0.212142	0.149584	0.148028
2024-04-19 16:15:00	0.230840	0.200828	0.169898	0.182518
2024-04-19 16:20:00	0.218375	0.213556	0.198984	0.200828

The model was validated on a portion of the dataset which was not used for training. A "train-test" split was performed. Additionally, the focus was on close price prediction only (lagged open values remained as features).

```
In [ ]: X = jnj_close[['Open Lag', 'Close Lag']] # Notice how the 'Open' column is implicitly dropped, unused.
y = jnj_close[['Close']]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.10, random_state=42)
X_train.head()
```

```
Out[ ]:
```

	Open Lag	Close Lag
Datetime		
2024-06-07 17:30:00	0.384340	0.350742
2024-05-07 20:50:00	0.462604	0.453513
2024-04-24 18:05:00	0.332410	0.315385
2024-04-22 17:30:00	0.495845	0.482741
2024-04-24 17:25:00	0.286242	0.305956

LSTMs require 3-dimensional inputs in the form of [samples, timesteps, features]. Inputs were reshaped.

```
In [ ]: X_train = X_train.values.reshape(X_train.shape[0], 1, X_train.shape[1])
X_test = X_test.values.reshape(X_test.shape[0], 1, X_test.shape[1])
```

An LSTM model was defined with multiple intermediate Dense layers. This increased both model accuracy and overall training time. It was then fitted.

```
In [ ]: model = Sequential()
model.add(LSTM(5, activation='relu', input_shape=(1,2)))
model.add(Dense(2))
model.add(Dense(2))
model.compile(optimizer='adam', loss='mse')
```

```
/home/stanis/Documents/programming/ATS_project/venv/lib64/python3.11/site-packages/keras/src/layers/rnn/rnn.py:
204: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, pr
efer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(**kwargs)
```

```
In [ ]: fitting_history = model.fit(X_train, y_train, epochs=15, batch_size=3, verbose=2)
```

```

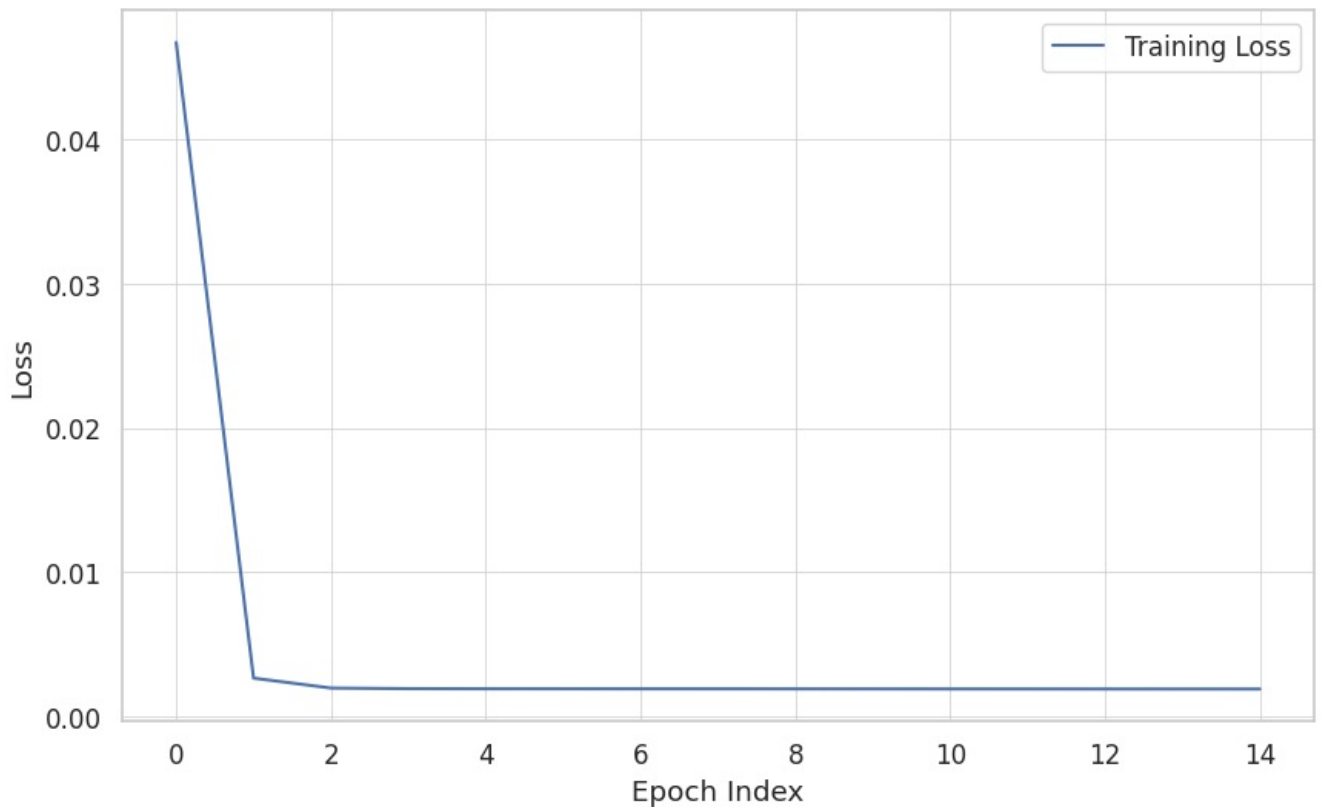
Epoch 1/15
771/771 - 4s - 5ms/step - loss: 0.0467
Epoch 2/15
771/771 - 2s - 2ms/step - loss: 0.0027
Epoch 3/15
771/771 - 2s - 2ms/step - loss: 0.0020
Epoch 4/15
771/771 - 2s - 2ms/step - loss: 0.0019
Epoch 5/15
771/771 - 2s - 2ms/step - loss: 0.0019
Epoch 6/15
771/771 - 2s - 2ms/step - loss: 0.0019
Epoch 7/15
771/771 - 2s - 2ms/step - loss: 0.0019
Epoch 8/15
771/771 - 2s - 2ms/step - loss: 0.0019
Epoch 9/15
771/771 - 2s - 2ms/step - loss: 0.0019
Epoch 10/15
771/771 - 2s - 2ms/step - loss: 0.0019
Epoch 11/15
771/771 - 2s - 2ms/step - loss: 0.0019
Epoch 12/15
771/771 - 2s - 2ms/step - loss: 0.0019
Epoch 13/15
771/771 - 2s - 2ms/step - loss: 0.0019
Epoch 14/15
771/771 - 2s - 2ms/step - loss: 0.0019
Epoch 15/15
771/771 - 2s - 2ms/step - loss: 0.0019

```

```

In [ ]: _, ax = plt.subplots()
ax.plot(fitting_history.history['loss'], label='Training Loss')
ax.set_xlabel('Epoch Index')
ax.set_ylabel('Loss')
plt.legend()
plt.show()

```



The model fitting history shows that the loss quickly diminished and remained constant after the initial epochs.

```

In [ ]: testPredict = model.predict(X_test)[: ,1]
testPredict = testPredict.reshape(len(testPredict), 1)

trainPredict = model.predict(X_train)[: ,1]
trainPredict = trainPredict.reshape(len(trainPredict), 1)

```

```

9/9 ————— 0s 28ms/step
73/73 ————— 0s 2ms/step

```

```

In [ ]: trainScore = np.sqrt(mean_squared_error(y_train, trainPredict))
print('Train Score: %.4f RMSE' % (trainScore))

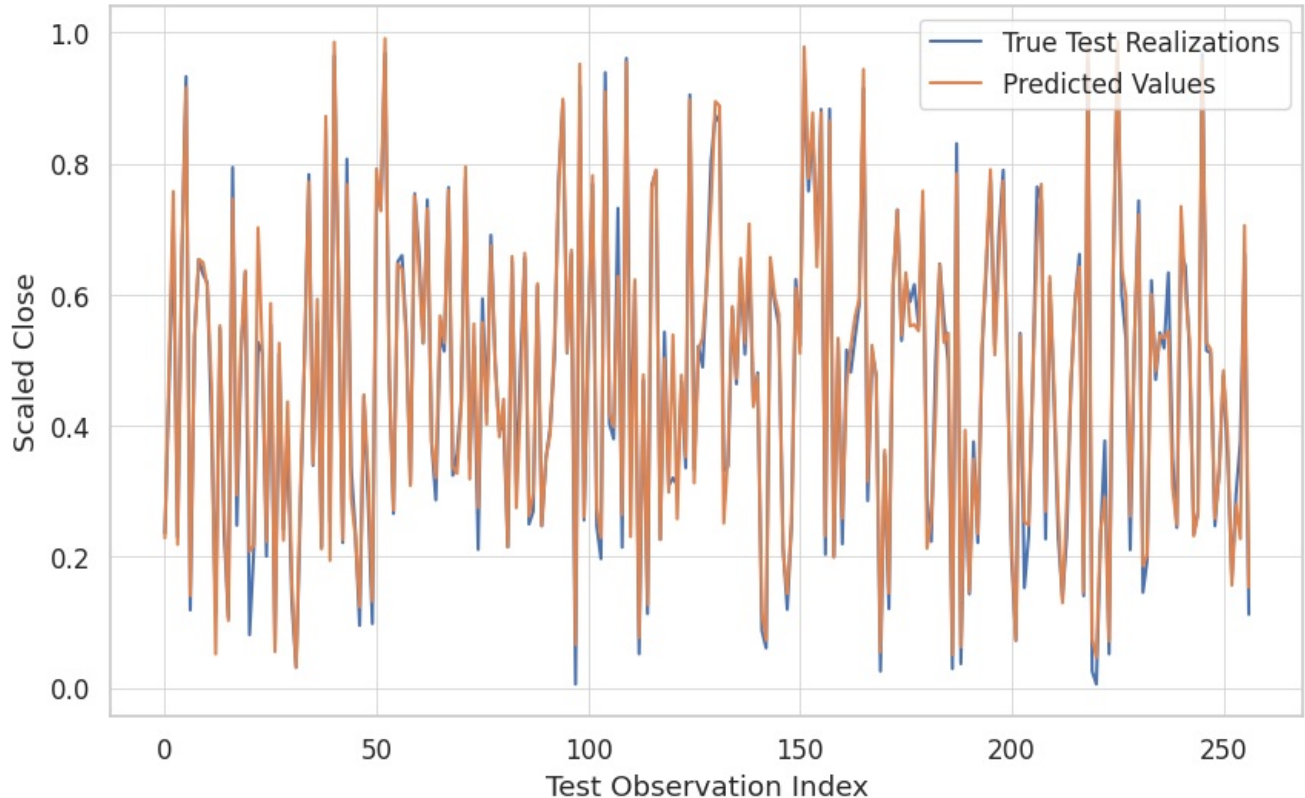
```

```
testScore = np.sqrt(mean_squared_error(y_test, testPredict))
print('Test Score: %.4f RMSE' % (testScore))
```

Train Score: 0.0435 RMSE

Test Score: 0.0371 RMSE

```
In [ ]: _, ax = plt.subplots()
ax.plot(y_test.values, label='True Test Realizations')
ax.plot(testPredict, label='Predicted Values')
ax.set_xlabel('Test Observation Index')
ax.set_ylabel('Scaled Close')
plt.legend()
plt.show()
```



As seen on the figure above, predictions of the scaled close price were very accurate. The test score was only slightly worse than the train score (as expected), which meant that the procedure did not overfit the model.

## Strategy & Results

The strategy was to set a long position when the model predicted a positive return. It did nothing if the predicted return was negative to limit trading frequency.

Firstly, the test sample data was arranged into a data frame, because LSTM forced data into numpy arrays above. The below dataframe still operated on scaled (to [0,1]) values.

```
In [ ]: lstm_strategy = pd.DataFrame(data=[y_test.values.T[0], testPredict.T[0]].transpose()
lstm_strategy.columns = columns=['True', 'Predicted']
lstm_strategy.head()
```

```
Out[ ]:   True Predicted
0  0.236232  0.228875
1  0.461055  0.493040
2  0.701483  0.757850
3  0.230056  0.218773
4  0.653398  0.653928
```

The scaling applied on the dataset was inverted to get the actual values of the close price.

```
In [ ]: predicted = testPredict.T[0].reshape(-1, 1)
predicted = np.c_[ predicted, predicted, predicted, predicted ] # The original transform had 4 features, we only
predicted = scaler.inverse_transform(predicted)[:,-1]

actual = y_test.values.T[0].reshape(-1, 1)
```

```
actual = np.c_[actual, actual, actual, actual]
actual = scaler.inverse_transform(actual)[: ,1]
```

```
In [ ]: lstm_strategy = pd.DataFrame(data=[actual, predicted]).transpose()
lstm_strategy.columns = columns=['True', 'Predicted']
lstm_strategy.head()
```

```
Out[ ]:
```

	True	Predicted
0	146.5555	146.477478
1	148.9400	149.279236
2	151.4900	152.087830
3	146.4900	146.370331
4	150.9800	150.985626

Predicted and true returns on the out-of-sample test period were calculated.

```
In [ ]: lstm_strategy['Predicted Return'] = (lstm_strategy['Predicted'] - lstm_strategy['Predicted'].shift(1)).shift(-1)
lstm_strategy['True Return'] = (lstm_strategy['True'] - lstm_strategy['True'].shift(1)).shift(-1)

lstm_strategy = lstm_strategy.dropna()
lstm_strategy.head()
```

```
Out[ ]:
```

	True	Predicted	Predicted Return	True Return
0	146.5555	146.477478	2.801758	2.3845
1	148.9400	149.279236	2.808594	2.5500
2	151.4900	152.087830	-5.717499	-5.0000
3	146.4900	146.370331	4.615295	4.4900
4	150.9800	150.985626	2.782455	2.9650

The signal was defined to equal 1 only if predicted one-day return was positive. The strategy did not trade otherwise.

```
In [ ]: lstm_strategy = lstm_strategy.assign(signal = lambda row: row['Predicted Return'] > 0)
lstm_strategy.head()
```

```
Out[ ]:
```

	True	Predicted	Predicted Return	True Return	signal
0	146.5555	146.477478	2.801758	2.3845	True
1	148.9400	149.279236	2.808594	2.5500	True
2	151.4900	152.087830	-5.717499	-5.0000	False
3	146.4900	146.370331	4.615295	4.4900	True
4	150.9800	150.985626	2.782455	2.9650	True

Strategy returns were defined to equal the actual returns obtained in the market that time less transaction costs defined above. It should be noted that the time index represented 5-minute observations. The evaluation period was therefore short.

```
In [ ]: lstm_strategy = lstm_strategy.assign(strategy_gross_returns = lambda row: row['signal'] * row['True Return'])
lstm_strategy['Did Position Change'] = (lstm_strategy['signal'].shift(1) == False) & (lstm_strategy['signal'] !=
lstm_strategy = lstm_strategy.assign(strategy_net_returns = lambda row: row['strategy_gross_returns'] - row['Di
lstm_strategy['Cumulative Gross Returns'] = lstm_strategy.cumsum()['strategy_gross_returns']
lstm_strategy['Cumulative Net Returns'] = lstm_strategy.cumsum()['strategy_net_returns']
lstm_strategy.head(10)
```

Out[ ]:

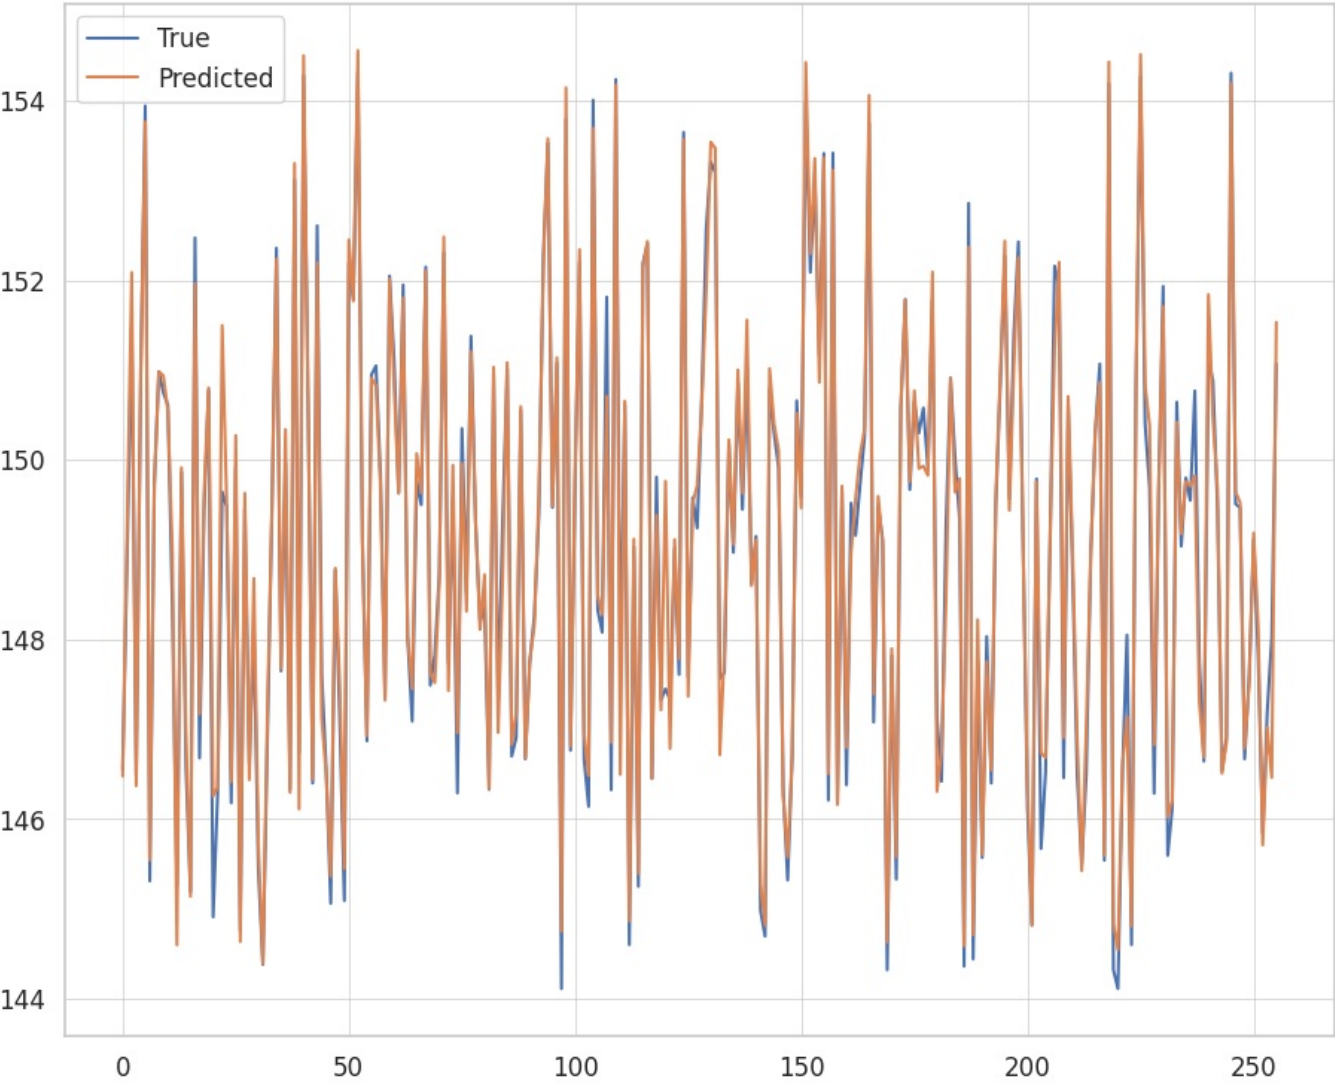
	True	Predicted	Predicted Return	True Return	signal	strategy_gross_returns	Did Position Change	strategy_net_returns	Cumulative Gross Returns	Cumulative Net Returns
0	146.5555	146.477478	2.801758	2.3845	True	2.3845	False	2.384500	2.3845	2.384500
1	148.9400	149.279236	2.808594	2.5500	True	2.5500	False	2.550000	4.9345	4.934500
2	151.4900	152.087830	-5.717499	-5.0000	False	-0.0000	False	-0.000000	4.9345	4.934500
3	146.4900	146.370331	4.615295	4.4900	True	4.4900	True	4.416755	9.4245	9.351255
4	150.9800	150.985626	2.782455	2.9650	True	2.9650	False	2.965000	12.3895	12.316255
5	153.9450	153.768082	-8.224365	-8.6350	False	-0.0000	False	-0.000000	12.3895	12.316255
6	145.3100	145.543716	4.159760	4.4300	True	4.4300	True	4.357345	16.8195	16.673600
7	149.7400	149.703476	1.281158	1.2450	True	1.2450	False	1.245000	18.0645	17.918600
8	150.9850	150.984634	-0.047119	-0.2350	False	-0.0000	False	-0.000000	18.0645	17.918600
9	150.7500	150.937515	-0.393753	-0.1400	False	-0.0000	False	-0.000000	18.0645	17.918600

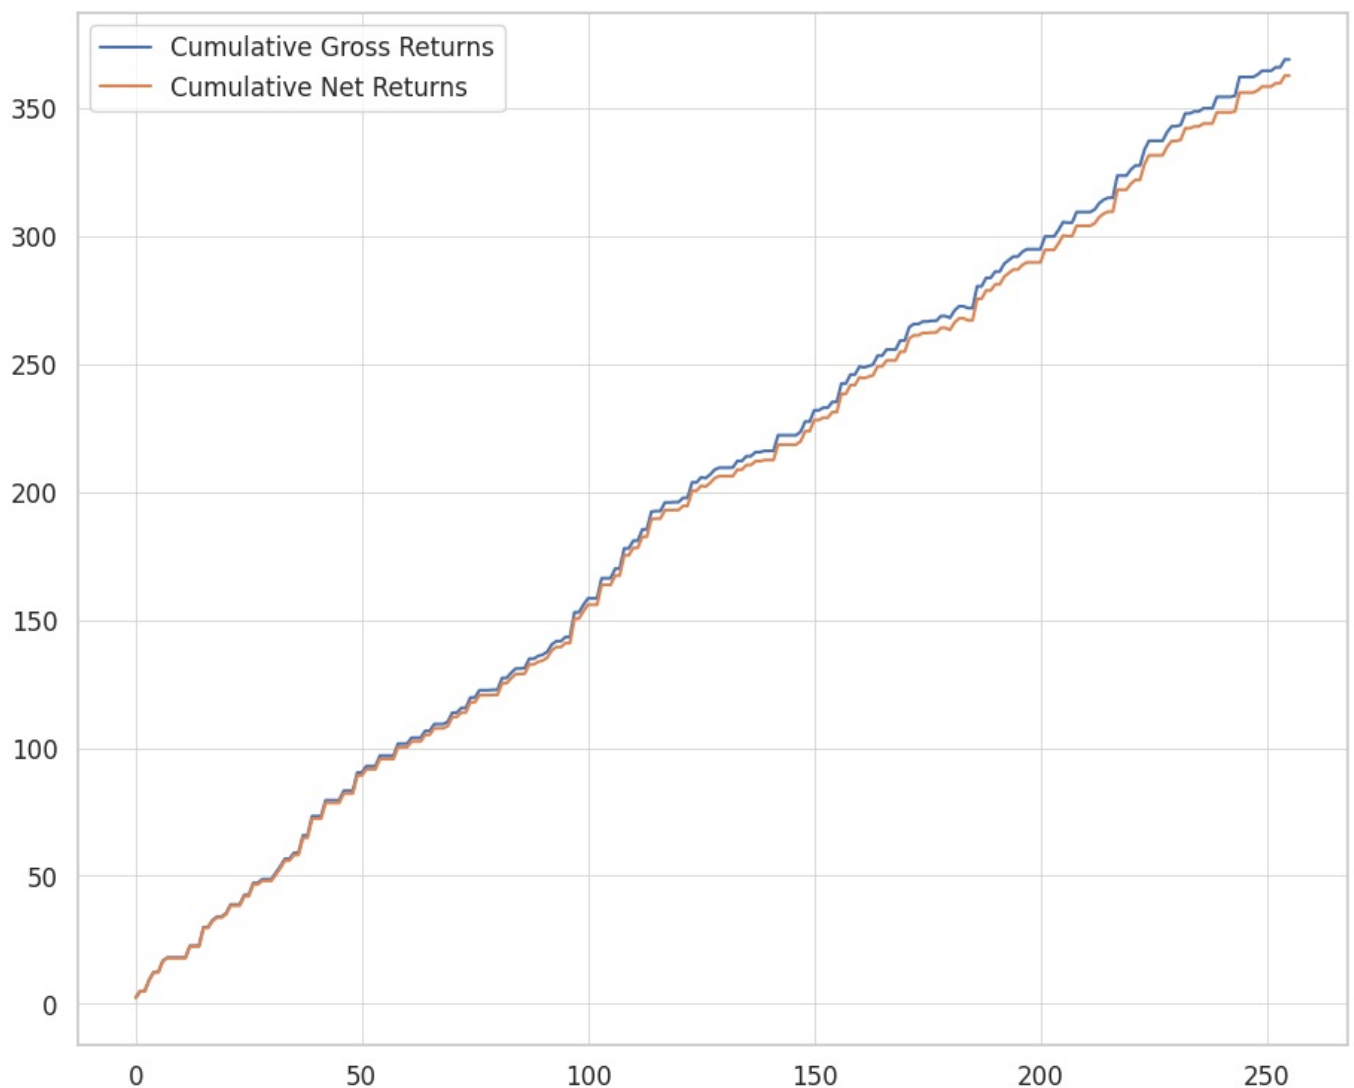
Strategy results were plotted against true and predicted returns. They included transaction costs defined above.

In [ ]:

```
lstm_strategy[['True', 'Predicted']].plot(figsize=(11,9))
lstm_strategy[['Cumulative Gross Returns', 'Cumulative Net Returns']].plot(figsize=(11,9))
```

Out[ ]: <Axes: >





The LSTM model was assessed using the total Sharpe Ratio on the out-of-sample test period. Simple statistics on the number of performed trades were calculated.

```
In [ ]: lstm_number_of_trades = lstm_strategy.cumsum()['signal'].iloc[-1]
lstm_proportion_of_trades = lstm_number_of_trades / len(lstm_strategy)
print(f"LSTM model made {lstm_number_of_trades} trades, it traded {lstm_proportion_of_trades:.{2}}% of times on
```

LSTM model made 130 trades, it traded 0.51% of times on average.

```
In [ ]: lstm_sharpe = qs.stats.sharpe(lstm_strategy[['Cumulative Net Returns']])[0]
print(f"LSTM model total Net Sharpe Ratio is equal to {lstm_sharpe:.{3}}")
```

LSTM model total Net Sharpe Ratio is equal to 3.79

```
/tmp/ipykernel_6274/3507800921.py:1: FutureWarning: Series.__getitem__ treating keys as positions is deprecated.
In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
lstm_sharpe = qs.stats.sharpe(lstm_strategy[['Cumulative Net Returns']])[0]
```

The Sharpe ratio was very high. Even including transaction costs, the strategy generated profits steadily and, because it traded only long, the variance of cumulative returns was very low.

## Conclusion

The underlying series of Johnson & Johnson equity returns displayed leptokurtic and autoregressive properties across the 19/04/2024-07/06/2024 period, assuming 5-minute prices.



The AR(9)-GARCH(1,2) model produced poor results despite the confirmed persistence of returns. The GARCH component was not able to forecast volatility sufficiently which triggered too many costly trades. The reaction to the signals was inappropriate for the frequency of data used - a mean-reversion strategy and a longer interval between observations may have been beneficial at increasing the -17.5 Sharpe ratio.

The LSTM model performed very well with an impressive Sharpe ratio of 3.79. The high result was attributed to the specification of the LSTM model, which captured short relationships in time very well. The diminishing gradients problem did not affect the prediction adversely. Additionally, the forecasting period of the stock was very stable and behaved like a mean-reverting process. Such a relationship was easily discovered by the model with a dense neural layer, as it effectively partially reduced to a scalar multiplication by -1. The model results could have been worse during a more volatile and trending period on the market. Nonetheless, the model results were impressive and were a good choice for a calm market.

## References

- [https://en.wikipedia.org/wiki/Autoregressive\\_conditional\\_heteroskedasticity](https://en.wikipedia.org/wiki/Autoregressive_conditional_heteroskedasticity)
- <https://arch.readthedocs.io/en/latest/univariate/introduction.html>
- <https://www.statsmodels.org/stable/index.html>
- <https://machinelearningmastery.com/multivariate-time-series-forecasting-lstms-keras/>
- <https://stackoverflow.com/questions/8486294/how-do-i-add-an-extra-column-to-a-numpy-array>
- <https://pypi.org/project/QuantStats/>
- <https://stackoverflow.com/questions/57216718/how-to-inverse-transform-the-predicted-values-in-a-multivariate-time-series-lstm>
- [https://en.wikipedia.org/wiki/Long\\_short-term\\_memory](https://en.wikipedia.org/wiki/Long_short-term_memory)

## Distribution of Work

Distribution of work matches our initial division of work from the project proposal. In general:

Adam:

- GitHub repo setup
- Definition of classes used in the first trading strategy
- Descriptive statistics of the dataset
- Implementation and testing of AR-GARCH
- Evaluation of the AR-GARCH, plotting results and strategy returns

Maciek:

- Data download
- Reading and preparing the time series
- Implementation and testing of the LSTM model
- Evaluation of LSTM, plotting results and strategy returns